

In [30]:

```
# import needed libraries
import pandas as pd
import numpy as np
```

## Task1 (1)

Load the row data into Pandas with the meta information being removed.

In [31]:

```
# read data from csv file
df = pd.read_csv('ckd-dataset-v2.csv')
df = df.drop([0,1]) #remove descriptive row and null row (remove the meta information)
```

## View the data

In [32]:

```
df.head()#view data
```

Out[32]:

|   | bp<br>(Diastolic) | bp<br>limit | sg                  | al          | class | rbc | su     | pc | pcc | ba | ... | htn | dm | cad | appet | pe | ane |
|---|-------------------|-------------|---------------------|-------------|-------|-----|--------|----|-----|----|-----|-----|----|-----|-------|----|-----|
| 2 | 0                 | 0           | 1.019<br>-<br>1.021 | 1<br>-<br>1 | ckd   | 0   | <<br>0 | 0  | 0   | 0  | ... | 0   | 0  | 0   | 0     | 0  | 0   |
| 3 | 0                 | 0           | 1.009<br>-<br>1.011 | <<br>0      | ckd   | 0   | <<br>0 | 0  | 0   | 0  | ... | 0   | 0  | 0   | 0     | 0  | 0   |
| 4 | 0                 | 0           | 1.009<br>-<br>1.011 | ≥<br>4      | ckd   | 1   | <<br>0 | 1  | 0   | 1  | ... | 0   | 0  | 0   | 1     | 0  | 0   |
| 5 | 1                 | 1           | 1.009<br>-<br>1.011 | 3<br>-<br>3 | ckd   | 0   | <<br>0 | 0  | 0   | 0  | ... | 0   | 0  | 0   | 0     | 0  | 0   |
| 6 | 0                 | 0           | 1.015<br>-<br>1.017 | <<br>0      | ckd   | 0   | <<br>0 | 0  | 0   | 0  | ... | 0   | 1  | 0   | 1     | 1  | 0   |

5 rows × 29 columns

## Task 1 (2)

Apply ordinal encoding to all columns contains numerical bins. For instance, in the “age” column, 0 for “<12”, 1 for “12-20”, 2 for “20-27”, etc.

## Printing all unique values for each row in the dataset

In [33]:

```
for names in df.columns: #print all unique values of each row to create encoding map
    print(names, end = ' -> ')
    print(df[names].unique())
```

```
bp (Diastolic) -> ['0' '1']
bp limit -> ['0' '1' '2']
sg -> ['1.019 - 1.021' '1.009 - 1.011' '1.015 - 1.017' '≥ 1.023' '< 1.007']
al -> ['1 - 1' '< 0' '≥ 4' '3 - 3' '2 - 2']
class -> ['ckd' 'notckd']
rbc -> ['0' '1']
su -> ['< 0' '4 - 4' '2 - 2' '3 - 4' '1 - 2' '≥ 4']
pc -> ['0' '1']
pcc -> ['0' '1']
ba -> ['0' '1']
bgr -> ['< 112' '112 - 154' '154 - 196' '406 - 448' '238 - 280' '196 - 238'
'≥ 448' '280 - 322' '364 - 406' '322 - 364']
bu -> ['< 48.1' '48.1 - 86.2' '200.5 - 238.6' '124.3 - 162.4' '86.2 - 124.3'
'162.4 - 200.5' '≥ 352.9' '238.6 - 276.7']
sod -> ['138 - 143' '133 - 138' '123 - 128' '143 - 148' '148 - 153' '< 118'
'128 - 133' '118 - 123' '≥ 158']
sc -> ['< 3.65' '3.65 - 6.8' '16.25 - 19.4' '6.8 - 9.95' '13.1 - 16.2 5'
'9.95 - 13.1' '≥ 28.85']
pot -> ['< 7.31' '≥ 42.59' '7.31 - 11.72' '38.18 - 42.59']
hemo -> ['11.3 - 12.6' '8.7 - 10' '13.9 - 15.2' '≥ 16.5' '10 - 11.3'
'7.4 - 8.7'
'12.6 - 13.9' '15.2 - 16.5' '< 6.1' '6.1 - 7.4']
pcv -> ['33.5 - 37.4' '29.6 - 33.5' '41.3 - 45.2' '37.4 - 41.3' '≥ 49.1'
'21.8 - 25.7' '45.2 - 49.1' '< 17.9' '25.7 - 29.6' '17.9 - 21.8']
rbcc -> ['4.46 - 5.05' '5.05 - 5.64' '3.28 - 3.87' '3.87 - 4.46' '6.23 - 6.82'
'5.64 - 6.23' '2.69 - 3.28' '< 2.69' '≥ 7.41']
wbcc -> ['7360 - 9740' '12120 - 14500' '14500 - 16880' '4980 - 7360'
'< 4980'
'9740 - 12120' '16880 - 19260' '≥ 24020' '19260 - 21640']
htn -> ['0' '1']
dm -> ['0' '1']
cad -> ['0' '1']
appet -> ['0' '1']
pe -> ['0' '1']
ane -> ['0' '1']
grf -> ['≥ 227.944' '127.281 - 152.446' '102.115 - 127.281' '177.612 - 202.778'
'26.6175 - 51.7832' '51.7832 - 76.949' '76.949 - 102.115'
'152.446 - 177.612' '202.778 - 227.944' '< 26.6175' ' p '']
stage -> ['s1' 's4' 's3' 's2' 's5']
affected -> ['1' '0']
age -> ['< 12' '12 - 20' '20 - 27' '27 - 35' '35 - 43' '43 - 51' '51 - 59'
'59 - 66' '66 - 74' '≥ 74']
```

Notice that grf consists of numerical bins except for 1 value which is 'p', therefore the row needs to be cleaned up.

In [34]:

```
df = df[df.grf != ' p ']#removing the row with p.
```

In [35]:

```

#running unique values for dataframe again.
for names in df.columns: #print all unique values of each row to create encoding map
    print(names, end = ' -> ')
    print(df[names].unique())# note that p is now missing from grf

bp (Diastolic) -> ['0' '1']
bp limit -> ['0' '1' '2']
sg -> ['1.019 - 1.021' '1.009 - 1.011' '1.015 - 1.017' '≥ 1.023' '< 1.007']
al -> ['1 - 1' '< 0' '≥ 4' '3 - 3' '2 - 2']
class -> ['ckd' 'notckd']
rbc -> ['0' '1']
su -> ['< 0' '4 - 4' '2 - 2' '3 - 4' '1 - 2' '≥ 4']
pc -> ['0' '1']
pcc -> ['0' '1']
ba -> ['0' '1']
bgr -> ['< 112' '112 - 154' '154 - 196' '406 - 448' '238 - 280' '196 - 238'
'≥ 448' '280 - 322' '364 - 406' '322 - 364']
bu -> ['< 48.1' '48.1 - 86.2' '200.5 - 238.6' '124.3 - 162.4' '86.2 - 124.3'
'162.4 - 200.5' '≥ 352.9' '238.6 - 276.7']
sod -> ['138 - 143' '133 - 138' '123 - 128' '143 - 148' '148 - 153' '< 118'
'128 - 133' '118 - 123' '≥ 158']
sc -> ['< 3.65' '3.65 - 6.8' '16.25 - 19.4' '6.8 - 9.95' '13.1 - 16.25'
'9.95 - 13.1' '≥ 28.85']
pot -> ['< 7.31' '≥ 42.59' '7.31 - 11.72' '38.18 - 42.59']
hemo -> ['11.3 - 12.6' '8.7 - 10' '13.9 - 15.2' '≥ 16.5' '10 - 11.3'
'7.4 - 8.7'
'12.6 - 13.9' '15.2 - 16.5' '< 6.1' '6.1 - 7.4']
pcv -> ['33.5 - 37.4' '29.6 - 33.5' '41.3 - 45.2' '37.4 - 41.3' '≥ 49.1'
'21.8 - 25.7' '45.2 - 49.1' '< 17.9' '25.7 - 29.6' '17.9 - 21.8']
rbcc -> ['4.46 - 5.05' '5.05 - 5.64' '3.28 - 3.87' '3.87 - 4.46' '6.23 - 6.82'
'5.64 - 6.23' '2.69 - 3.28' '< 2.69' '≥ 7.41']
wbcc -> ['7360 - 9740' '12120 - 14500' '14500 - 16880' '4980 - 7360'
'< 4980'
'9740 - 12120' '16880 - 19260' '≥ 24020' '19260 - 21640']
htn -> ['0' '1']
dm -> ['0' '1']
cad -> ['0' '1']
appet -> ['0' '1']
pe -> ['0' '1']
ane -> ['0' '1']
grf -> ['≥ 227.944' '127.281 - 152.446' '102.115 - 127.281' '177.612 - 202.778'
'26.6175 - 51.7832' '51.7832 - 76.949' '76.949 - 102.115'
'152.446 - 177.612' '202.778 - 227.944' '< 26.6175']
stage -> ['s1' 's4' 's3' 's2' 's5']
affected -> ['1' '0']
age -> ['< 12' '12 - 20' '20 - 27' '27 - 35' '35 - 43' '43 - 51' '51 - 59'
'59 - 66' '66 - 74' '≥ 74']

```

We can observe that columns sg,al,al,su,bgr,bu,sod,sc,pot,hemo,pcv,rbcc,wbcc,grf,age needs to be encoded.

**We created a map for each column in accordance to the numerical bin it represents to preserve the relationship between attributes**

In [36]:

```

#Apply ordinal encoding to all columns contains numerical bins. For instance, in the
#for "<12", 1 for "12-20", 2 for "20-27", etc.
encoding_map = {"age" : {'< 12':0, '12 - 20':1, '20 - 27':2, '27 - 35':3, '35 - 43':
    '51 - 59':6, '59 - 66':7, '66 - 74':8, '≥ 74':9},
    "grf" : {'≥ 227.944':9, '127.281 - 152.446':5, '102.115 - 127.281':4
:7, '26.6175 - 51.7832':1, '51.7832 - 76.949':2, '76.949 - 102.115':3,
    '152.446 - 177.612':6, '202.778 - 227.944':8, '< 26.6175':0},
    'wbcc': {'7360 - 9740':2, '12120 - 14500':4, '14500 - 16880':5, '4980
    '9740 - 12120':3, '16880 - 19260':6, '≥ 24020':8, '19260 - 21640':7},
    'rbcc': {'4.46 - 5.05':4, '5.05 - 5.64':5, '3.28 - 3.87':2, '3.87 - 4.
    '5.64 - 6.23':6, '2.69 - 3.28':1, '< 2.69':0, '≥ 7.41':8},
    'bgr' : {'< 112' : 0, '112 - 154' : 1, '154 - 196' : 2, '406 - 448' :
'bu' : {'< 48.1' : 0, '48.1 - 86.2' : 1, '200.5 - 238.6' : 5, '124.3 - 162.4' : 3,
'sod' : {'138 - 143' : 5, '133 - 138' : 4, '123 - 128' : 2, '143 - 148' : 6, '148 -
'sc' : {'< 3.65' : 0, '3.65 - 6.8' : 1, '16.25 - 19.4' : 5, '6.8 - 9.95' : 2, '13.1
'pot' : {'< 7.31' : 0, '≥ 42.59' : 3, '7.31 - 11.72' : 1, '38.18 - 42.59' : 2},
'hemo' : {'11.3 - 12.6' : 5, '8.7 - 10' : 3, '13.9 - 15.2' : 7, '≥ 16.5' : 9, '10 -
'pcv' : {'33.5 - 37.4' : 5, '29.6 - 33.5' : 4, '41.3 - 45.2' : 7, '37.4 - 41.3' : 6,
'su' : {'< 0' : 0, '4 - 4' : 4, '2 - 2' : 2, '3 - 4' : 3, '1 - 2' : 1, '≥ 4' : 5},
'sg' : {'1.019 - 1.021' : 3, '1.009 - 1.011' : 1, '1.015 - 1.017' : 2, '≥ 1.023' : 4
'al' : {'1 - 1' : 1, '< 0' : 0, '≥ 4' : 4, '3 - 3' : 3, '2 - 2' : 2}}
encoded_df = df.replace(encoding_map)
#running unique values for dataframe again.
for names in encoded_df.columns: #print all unique values of each row to verify enco
    print(names, end = ' -> ')
    print(encoded_df[names].unique()) # note that all numerical bins are encoded

```

```

bp (Diastolic) -> ['0' '1']
bp limit -> ['0' '1' '2']
sg -> [3 1 2 4 0]
al -> [1 0 4 3 2]
class -> ['ckd' 'notckd']
rbc -> ['0' '1']
su -> [0 4 2 3 1 5]
pc -> ['0' '1']
pcc -> ['0' '1']
ba -> ['0' '1']
bgr -> [0 1 2 8 4 3 9 5 7 6]
bu -> [0 1 5 3 2 4 7 6]
sod -> [5 4 2 6 7 0 3 1 8]
sc -> [0 1 5 2 4 3 6]
pot -> [0 3 1 2]
hemo -> [5 3 7 9 4 2 6 8 0 1]
pcv -> [5 4 7 6 9 2 8 0 3 1]
rbcc -> [4 5 2 3 7 6 1 0 8]
wbcc -> [2 4 5 1 0 3 6 8 7]
htn -> ['0' '1']
dm -> ['0' '1']
cad -> ['0' '1']
appet -> ['0' '1']
pe -> ['0' '1']
ane -> ['0' '1']
grf -> [9 5 4 7 1 2 3 6 8 0]
stage -> ['s1' 's4' 's3' 's2' 's5']
affected -> ['1' '0']
age -> [0 1 2 3 4 5 6 7 8 9]

```

## Task 1 (3)

Apply one-hot-encoding to the “stage” column.

In [37]:

```
#first view unique stages  
encoded_df["stage"].unique()
```

Out[37]:

```
array(['s1', 's4', 's3', 's2', 's5'], dtype=object)
```

We can observe that the only possible values for stage is s1,s2,s3 and s4

In [38]:

```
#using pd.get_dummies to one hot encode the stage column  
dummy = pd.get_dummies(encoded_df.stage,prefix='Stage')  
dummy.head()
```

Out[38]:

|   | Stage_s1 | Stage_s2 | Stage_s3 | Stage_s4 | Stage_s5 |
|---|----------|----------|----------|----------|----------|
| 2 | 1        | 0        | 0        | 0        | 0        |
| 3 | 1        | 0        | 0        | 0        | 0        |
| 4 | 1        | 0        | 0        | 0        | 0        |
| 5 | 1        | 0        | 0        | 0        | 0        |
| 6 | 1        | 0        | 0        | 0        | 0        |

Creating a dummy table for one hot encoding

In [39]:

```
#join dummy dataframe to original data frame
new_df = encoded_df.join(dummy)
new_df
```

Out[39]:

|     | bp<br>(Diastolic) | bp<br>limit | sg  | al  | class  | rbc | su  | pc  | pcc | ba  | ... | ane | grf | stage | affected | age |
|-----|-------------------|-------------|-----|-----|--------|-----|-----|-----|-----|-----|-----|-----|-----|-------|----------|-----|
| 2   | 0                 | 0           | 3   | 1   | ckd    | 0   | 0   | 0   | 0   | 0   | ... | 0   | 9   | s1    | 1        | 0   |
| 3   | 0                 | 0           | 1   | 0   | ckd    | 0   | 0   | 0   | 0   | 0   | ... | 0   | 9   | s1    | 1        | 0   |
| 4   | 0                 | 0           | 1   | 4   | ckd    | 1   | 0   | 1   | 0   | 1   | ... | 0   | 5   | s1    | 1        | 0   |
| 5   | 1                 | 1           | 1   | 3   | ckd    | 0   | 0   | 0   | 0   | 0   | ... | 0   | 5   | s1    | 1        | 0   |
| 6   | 0                 | 0           | 2   | 0   | ckd    | 0   | 0   | 0   | 0   | 0   | ... | 0   | 5   | s1    | 1        | 1   |
| ... | ...               | ...         | ... | ... | ...    | ... | ... | ... | ... | ... | ... | ... | ... | ...   | ...      | ... |
| 197 | 1                 | 2           | 3   | 0   | ckd    | 0   | 0   | 0   | 0   | 0   | ... | 1   | 1   | s3    | 1        | 9   |
| 198 | 0                 | 0           | 3   | 0   | ckd    | 0   | 0   | 0   | 0   | 0   | ... | 1   | 0   | s4    | 1        | 9   |
| 199 | 1                 | 1           | 4   | 0   | notckd | 0   | 0   | 0   | 0   | 0   | ... | 0   | 2   | s2    | 0        | 9   |
| 200 | 1                 | 1           | 4   | 0   | notckd | 0   | 0   | 0   | 0   | 0   | ... | 0   | 4   | s1    | 0        | 9   |
| 201 | 1                 | 1           | 1   | 2   | ckd    | 0   | 2   | 0   | 0   | 0   | ... | 0   | 0   | s4    | 1        | 9   |

199 rows × 34 columns

We can observe that the dummy table is now appended to the end of the data frame

In [40]:

```
new_df = new_df.drop(columns="stage") #dropping original stage column
```

Dropping the stage column from the dataframe as the information has been encoded in the new columns

## Task 1 (4)

Find out and rank the correlations between the “class” column and all other columns



In [41]:

```
#Get information of all the data types of the dataset
new_df.info()
#some binary values are still object data types as they are wrapped in '0' and '1'
#class column is also in ckd and notckd which needs to be encoded to 0 and 1
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 199 entries, 2 to 201
Data columns (total 33 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   bp (Diastolic)         199 non-null    object
 1   bp limit               199 non-null    object
 2   sg                    199 non-null    int64
 3   al                    199 non-null    int64
 4   class                 199 non-null    object
 5   rbc                   199 non-null    object
 6   su                    199 non-null    int64
 7   pc                    199 non-null    object
 8   pcc                   199 non-null    object
 9   ba                    199 non-null    object
10  bgr                   199 non-null    int64
11  bu                    199 non-null    int64
12  sod                   199 non-null    int64
13  sc                    199 non-null    int64
14  pot                   199 non-null    int64
15  hemo                  199 non-null    int64
16  pcv                   199 non-null    int64
17  rbcc                  199 non-null    int64
18  wbcc                  199 non-null    int64
19  htn                   199 non-null    object
20  dm                    199 non-null    object
21  cad                   199 non-null    object
22  appet                 199 non-null    object
23  pe                    199 non-null    object
24  ane                   199 non-null    object
25  grf                   199 non-null    int64
26  affected              199 non-null    object
27  age                   199 non-null    int64
28  Stage_s1              199 non-null    uint8
29  Stage_s2              199 non-null    uint8
30  Stage_s3              199 non-null    uint8
31  Stage_s4              199 non-null    uint8
32  Stage_s5              199 non-null    uint8
dtypes: int64(14), object(14), uint8(5)
memory usage: 54.2+ KB
```

As you can observe, the dataframe still consists of not numerical attributes in the object data type columns, therefore we need to convert the str values representing integers into true integers

In [42]:

```
encode_class = {"class":{"ckd":1,'notckd':0}}
new_df = new_df.replace(encode_class)#change ckd and notckd to 1 and 0 respectively
new_df
```

Out[42]:

|     | bp<br>(Diastolic) | bp<br>limit | sg  | al  | class | rbc | su  | pc  | pcc | ba  | ... | pe  | ane | grf | affected | age | Sta |
|-----|-------------------|-------------|-----|-----|-------|-----|-----|-----|-----|-----|-----|-----|-----|-----|----------|-----|-----|
| 2   | 0                 | 0           | 3   | 1   | 1     | 0   | 0   | 0   | 0   | 0   | ... | 0   | 0   | 9   | 1        | 0   |     |
| 3   | 0                 | 0           | 1   | 0   | 1     | 0   | 0   | 0   | 0   | 0   | ... | 0   | 0   | 9   | 1        | 0   |     |
| 4   | 0                 | 0           | 1   | 4   | 1     | 1   | 0   | 1   | 0   | 1   | ... | 0   | 0   | 5   | 1        | 0   |     |
| 5   | 1                 | 1           | 1   | 3   | 1     | 0   | 0   | 0   | 0   | 0   | ... | 0   | 0   | 5   | 1        | 0   |     |
| 6   | 0                 | 0           | 2   | 0   | 1     | 0   | 0   | 0   | 0   | 0   | ... | 1   | 0   | 5   | 1        | 1   |     |
| ... | ...               | ...         | ... | ... | ...   | ... | ... | ... | ... | ... | ... | ... | ... | ... | ...      | ... | ... |
| 197 | 1                 | 2           | 3   | 0   | 1     | 0   | 0   | 0   | 0   | 0   | ... | 0   | 1   | 1   | 1        | 9   |     |
| 198 | 0                 | 0           | 3   | 0   | 1     | 0   | 0   | 0   | 0   | 0   | ... | 0   | 1   | 0   | 1        | 9   |     |
| 199 | 1                 | 1           | 4   | 0   | 0     | 0   | 0   | 0   | 0   | 0   | ... | 0   | 0   | 2   | 0        | 9   |     |
| 200 | 1                 | 1           | 4   | 0   | 0     | 0   | 0   | 0   | 0   | 0   | ... | 0   | 0   | 4   | 0        | 9   |     |
| 201 | 1                 | 1           | 1   | 2   | 1     | 0   | 2   | 0   | 0   | 0   | ... | 0   | 0   | 0   | 1        | 9   |     |

199 rows × 33 columns

Converting the class attribute to 1 or 0

In [43]:

```
new_df = new_df.astype(int)
new_df.info()#converting all strings numericals to string
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 199 entries, 2 to 201
Data columns (total 33 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   bp (Diastolic)        199 non-null    int64
 1   bp limit              199 non-null    int64
 2   sg                   199 non-null    int64
 3   al                   199 non-null    int64
 4   class                199 non-null    int64
 5   rbc                  199 non-null    int64
 6   su                   199 non-null    int64
 7   pc                   199 non-null    int64
 8   pcc                  199 non-null    int64
 9   ba                   199 non-null    int64
10  bgr                  199 non-null    int64
11  bu                   199 non-null    int64
12  sod                  199 non-null    int64
13  sc                   199 non-null    int64
14  pot                  199 non-null    int64
15  hemo                 199 non-null    int64
16  pcv                  199 non-null    int64
17  rbcc                 199 non-null    int64
18  wbcc                 199 non-null    int64
19  htn                  199 non-null    int64
20  dm                   199 non-null    int64
21  cad                  199 non-null    int64
22  appet               199 non-null    int64
23  pe                   199 non-null    int64
24  ane                  199 non-null    int64
25  grf                  199 non-null    int64
26  affected             199 non-null    int64
27  age                  199 non-null    int64
28  Stage_s1             199 non-null    int64
29  Stage_s2             199 non-null    int64
30  Stage_s3             199 non-null    int64
31  Stage_s4             199 non-null    int64
32  Stage_s5             199 non-null    int64
dtypes: int64(33)
memory usage: 61.0 KB
```

Converting all strings to ints

In [44]:

```
cormat = new_df.corr()  
cormat = cormat.abs()  
cormat['class'].sort_values(ascending=False)  
#sort the class correlation matrix after transforming p value to absolute value  
#ranking correlation by ascending order
```

Out[44]:

|                |          |
|----------------|----------|
| affected       | 1.000000 |
| class          | 1.000000 |
| hemo           | 0.765089 |
| pcv            | 0.699531 |
| sg             | 0.657702 |
| grf            | 0.637083 |
| rbcc           | 0.606591 |
| htn            | 0.604532 |
| Stage_s1       | 0.598828 |
| al             | 0.558429 |
| dm             | 0.548551 |
| sod            | 0.495498 |
| bu             | 0.424668 |
| pc             | 0.407015 |
| appet          | 0.377656 |
| bgr            | 0.366404 |
| pe             | 0.347838 |
| Stage_s5       | 0.341792 |
| ane            | 0.329595 |
| Stage_s3       | 0.323438 |
| Stage_s4       | 0.307020 |
| sc             | 0.301996 |
| pcc            | 0.298320 |
| su             | 0.286902 |
| rbc            | 0.285404 |
| Stage_s2       | 0.283929 |
| bp limit       | 0.273247 |
| cad            | 0.265454 |
| wbcc           | 0.235190 |
| age            | 0.232851 |
| ba             | 0.182130 |
| bp (Diastolic) | 0.098866 |
| pot            | 0.086149 |

Name: class, dtype: float64

**As we can observe from the correlation matrix, affected is exactly correlated to the class column with a p value of 1. Other notable values such as hemo, pcv, sg, grf and rbcc have medium to strong correlation to the class attribute, hinting that those features are good indicators of predicting Chronic Kidney Disease**