

- Kendrick Kee
- 7366814

Import Statements

In [1]:

```
import numpy as np
import pandas as pd
import csv
import matplotlib.pyplot as plt
import seaborn as sns
eps = np.finfo(float).eps
from numpy import log2 as log
import sys
import copy
import random
from pprint import pprint
```

In [2]:

```
df = pd.read_csv('secondary_data.csv', sep = ';')
df.head()
```

Out[2]:

	class	cap- diameter	cap- shape	cap- surface	cap- color	does- bruise- or- bleed	gill- attachment	gill- spacing	gill- color	stem- height	...	stern rock
0	p	15.26	x	g	o	f	e	NaN	w	16.95	...	
1	p	16.60	x	g	o	f	e	NaN	w	17.99	...	
2	p	14.07	x	g	o	f	e	NaN	w	17.80	...	
3	p	14.17	f	h	e	f	e	NaN	w	15.77	...	
4	p	14.64	x	h	o	f	e	NaN	w	16.53	...	

5 rows × 21 columns

Findings for current dataset.

Every column aside from the class column will consist of empty/null values. Therefore, dropping rows with empty columns is not a viable strategy in preparing this dataset for ML algorithms. I therefore suggest hot encoding all categorical columns and for binary classified columns such as "has-ring", "does-bruise-or-bleed" and "veil-type" (as the 2 possible values are either 'u' or null) into binary bins of 1 and 0

In [3]:

```
#encode the 3 binary datatype columns
encodingmap = {"class":{"p":1,'e':0},
               "has-ring":{"t":1,'f':0}}
df = df.replace(encodingmap)
df.head()
```

Out[3]:

	class	cap-diameter	cap-shape	cap-surface	cap-color	does-bruise-or-bleed	gill-attachment	gill-spacing	gill-color	stem-height	stem-width	stem-root
0	1	15.26	x	g	o	f	e	NaN	w	16.95
1	1	16.60	x	g	o	f	e	NaN	w	17.99
2	1	14.07	x	g	o	f	e	NaN	w	17.80
3	1	14.17	f	h	e	f	e	NaN	w	15.77
4	1	14.64	x	h	o	f	e	NaN	w	16.53

5 rows × 21 columns

Find columns with large amounts of empty data

In [4]:

```
df.isnull().sum()
```

Out[4]:

```
class                0
cap-diameter         0
cap-shape            0
cap-surface        14120
cap-color            0
does-bruise-or-bleed 0
gill-attachment      9884
gill-spacing        25063
gill-color           0
stem-height          0
stem-width           0
stem-root           51538
stem-surface        38124
stem-color           0
veil-type           57892
veil-color          53656
has-ring             0
ring-type            2471
spore-print-color    54715
habitat              0
season              0
dtype: int64
```

Here we can observe that only "gill-attachment" and "ring-type" has a reasonable amount of null values to be replaced while the other columns has null values that consist of more than half the entire row count.

Therefore, I will be dropping them as they will not be good predictors of the class type as there is insufficient information provided.

In [5]:

```
#dropping all columns with high null counts.
for col_names in list(df.columns):
    if df[col_names].isnull().sum() > (61069*0.2): #if column has more than 20% miss
        df = df.drop([col_names],axis=1)
df.head()
```

Out[5]:

	class	cap- diameter	cap- shape	cap- color	does- bruise- or- bleed	gill- attachment	gill- color	stem- height	stem- width	stem- color	has- ring	ring- type
0	1	15.26	x	o	f	e	w	16.95	17.09	w	1	g
1	1	16.60	x	o	f	e	w	17.99	18.19	w	1	g
2	1	14.07	x	o	f	e	w	17.80	17.74	w	1	g
3	1	14.17	f	e	f	e	w	15.77	15.98	w	1	p
4	1	14.64	x	o	f	e	w	16.53	17.20	w	1	p

In [6]:

```
#get missing values count again
df.isnull().sum()
```

Out[6]:

```
class                0
cap-diameter         0
cap-shape            0
cap-color            0
does-bruise-or-bleed 0
gill-attachment     9884
gill-color           0
stem-height          0
stem-width           0
stem-color           0
has-ring             0
ring-type           2471
habitat              0
season               0
dtype: int64
```

In [7]:

```
#fill the missing values with the mode of the column
for colname in ['gill-attachment', 'ring-type']:
    temp = df[colname].mode()[0]
    df[colname].fillna(temp, inplace=True)
df.isnull().sum()#get the null count again
```

Out[7]:

```
class                0
cap-diameter         0
cap-shape            0
cap-color            0
does-bruise-or-bleed 0
gill-attachment      0
gill-color           0
stem-height          0
stem-width           0
stem-color           0
has-ring             0
ring-type            0
habitat              0
season               0
dtype: int64
```

Findings

We can now observe that the dataset no longer has null values

In [8]:

```
#view object types to check if further encoding is needed
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 61069 entries, 0 to 61068
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   class                 61069 non-null  int64
1   cap-diameter          61069 non-null  float64
2   cap-shape             61069 non-null  object
3   cap-color             61069 non-null  object
4   does-bruise-or-bleed  61069 non-null  object
5   gill-attachment       61069 non-null  object
6   gill-color            61069 non-null  object
7   stem-height           61069 non-null  float64
8   stem-width            61069 non-null  float64
9   stem-color            61069 non-null  object
10  has-ring              61069 non-null  int64
11  ring-type             61069 non-null  object
12  habitat               61069 non-null  object
13  season                61069 non-null  object
dtypes: float64(3), int64(2), object(9)
memory usage: 6.5+ MB
```

In [9]:

```
#apply onehot encoding for categorical data which for this case would be columns with
#print(pd.get_dummies(df["cap-shape"],prefix='cap-shape'))
for key,value in dict(df.dtypes).items():
    if value == 'object':
        dummy = pd.get_dummies(df[key],prefix=key)
        df = df.join(dummy)
        df = df.drop(columns = key)

df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 61069 entries, 0 to 61068
```

```
Data columns (total 78 columns):
```

#	Column	Non-Null Count	Dtype
0	class	61069 non-null	int64
1	cap-diameter	61069 non-null	float64
2	stem-height	61069 non-null	float64
3	stem-width	61069 non-null	float64
4	has-ring	61069 non-null	int64
5	cap-shape_b	61069 non-null	uint8
6	cap-shape_c	61069 non-null	uint8
7	cap-shape_f	61069 non-null	uint8
8	cap-shape_o	61069 non-null	uint8
9	cap-shape_p	61069 non-null	uint8
10	cap-shape_s	61069 non-null	uint8
11	cap-shape_x	61069 non-null	uint8
12	cap-color_b	61069 non-null	uint8
13	cap-color_e	61069 non-null	uint8
14	cap-color_g	61069 non-null	uint8
15	cap-color_k	61069 non-null	uint8
16	cap-color_l	61069 non-null	uint8
17	cap-color_n	61069 non-null	uint8
18	cap-color_o	61069 non-null	uint8
19	cap-color_p	61069 non-null	uint8
20	cap-color_r	61069 non-null	uint8
21	cap-color_u	61069 non-null	uint8
22	cap-color_w	61069 non-null	uint8
23	cap-color_y	61069 non-null	uint8
24	does-bruise-or-bleed_f	61069 non-null	uint8
25	does-bruise-or-bleed_t	61069 non-null	uint8
26	gill-attachment_a	61069 non-null	uint8
27	gill-attachment_d	61069 non-null	uint8
28	gill-attachment_e	61069 non-null	uint8
29	gill-attachment_f	61069 non-null	uint8
30	gill-attachment_p	61069 non-null	uint8
31	gill-attachment_s	61069 non-null	uint8
32	gill-attachment_x	61069 non-null	uint8
33	gill-color_b	61069 non-null	uint8
34	gill-color_e	61069 non-null	uint8
35	gill-color_f	61069 non-null	uint8
36	gill-color_g	61069 non-null	uint8
37	gill-color_k	61069 non-null	uint8
38	gill-color_n	61069 non-null	uint8
39	gill-color_o	61069 non-null	uint8
40	gill-color_p	61069 non-null	uint8
41	gill-color_r	61069 non-null	uint8
42	gill-color_u	61069 non-null	uint8
43	gill-color_w	61069 non-null	uint8

```

44  gill-color_y          61069 non-null  uint8
45  stem-color_b          61069 non-null  uint8
46  stem-color_e          61069 non-null  uint8
47  stem-color_f          61069 non-null  uint8
48  stem-color_g          61069 non-null  uint8
49  stem-color_k          61069 non-null  uint8
50  stem-color_l          61069 non-null  uint8
51  stem-color_n          61069 non-null  uint8
52  stem-color_o          61069 non-null  uint8
53  stem-color_p          61069 non-null  uint8
54  stem-color_r          61069 non-null  uint8
55  stem-color_u          61069 non-null  uint8
56  stem-color_w          61069 non-null  uint8
57  stem-color_y          61069 non-null  uint8
58  ring-type_e           61069 non-null  uint8
59  ring-type_f           61069 non-null  uint8
60  ring-type_g           61069 non-null  uint8
61  ring-type_l           61069 non-null  uint8
62  ring-type_m           61069 non-null  uint8
63  ring-type_p           61069 non-null  uint8
64  ring-type_r           61069 non-null  uint8
65  ring-type_z           61069 non-null  uint8
66  habitat_d             61069 non-null  uint8
67  habitat_g             61069 non-null  uint8
68  habitat_h             61069 non-null  uint8
69  habitat_l             61069 non-null  uint8
70  habitat_m             61069 non-null  uint8
71  habitat_p             61069 non-null  uint8
72  habitat_u             61069 non-null  uint8
73  habitat_w             61069 non-null  uint8
74  season_a              61069 non-null  uint8
75  season_s              61069 non-null  uint8
76  season_u              61069 non-null  uint8
77  season_w              61069 non-null  uint8
dtypes: float64(3), int64(2), uint8(73)
memory usage: 6.6 MB

```

As observed, all columns are all now numerically represented.

After applying one hot encoding on the object rows, we now no longer have object data types

Now I will apply correlation to the class to view any significant features to use for model training as to improve model's accuracy and reduce the feature space and computational complexity.

In [10]:

```

pre_processed_df = df
correlation_matrix = df.corr()
for col_name, p_value in dict(correlation_matrix["class"].sort_values()).items():
    if p_value < 0.1 and p_value > -0.1: #selects features with a significant amount
        pre_processed_df = pre_processed_df.drop(columns=col_name) #drop columns of r

c = pre_processed_df.pop('class')
pre_processed_df.insert(18, 'class', c)
pre_processed_df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 61069 entries, 0 to 61068
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   cap-diameter           61069 non-null  float64
1   stem-height            61069 non-null  float64
2   stem-width             61069 non-null  float64
3   cap-shape_b            61069 non-null  uint8
4   cap-shape_o            61069 non-null  uint8
5   cap-color_b            61069 non-null  uint8
6   cap-color_e            61069 non-null  uint8
7   cap-color_n            61069 non-null  uint8
8   cap-color_r            61069 non-null  uint8
9   gill-attachment_a      61069 non-null  uint8
10  gill-attachment_e      61069 non-null  uint8
11  gill-attachment_p      61069 non-null  uint8
12  gill-color_n           61069 non-null  uint8
13  gill-color_w           61069 non-null  uint8
14  stem-color_f           61069 non-null  uint8
15  stem-color_w           61069 non-null  uint8
16  ring-type_z            61069 non-null  uint8
17  habitat_g              61069 non-null  uint8
18  class                  61069 non-null  int64
dtypes: float64(3), int64(1), uint8(15)
memory usage: 2.7 MB

```

In [11]:

```
pre_processed_df.describe()
```

Out[11]:

	cap-diameter	stem-height	stem-width	cap-shape_b	cap-shape_o	cap-color_b	c:
count	61069.000000	61069.000000	61069.000000	61069.000000	61069.000000	61069.000000	61069.000000
mean	6.733854	6.581538	12.149410	0.093239	0.056657	0.020141	0.000000
std	5.264845	3.370017	10.035955	0.290769	0.231188	0.140484	0.000000
min	0.380000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	3.480000	4.640000	5.210000	0.000000	0.000000	0.000000	0.000000
50%	5.860000	5.950000	10.190000	0.000000	0.000000	0.000000	0.000000
75%	8.540000	7.740000	16.570000	0.000000	0.000000	0.000000	0.000000
max	62.340000	33.920000	103.910000	1.000000	1.000000	1.000000	0.000000

As observed the feature space is not reduced to 18 columns and is now ready for training

Classification Helper Functions for Decision Tree Algorithm

Classification with Info gain and gini split helper functions

In [14]:

```
# Calculate Entropy of dataset
def findEntropy(df):
    Class = df.keys()[-1]
    entropy = 0
    values = df[Class].unique()

    for value in values:
        fraction = df[Class].value_counts()[value] / len(df[Class])
        entropy += -fraction * np.log2(fraction)
    return entropy
```


In [15]:

```
# Calculate Entropy by attribute
def findEntropyAttribute(df, attribute):
    Class = df.keys()[-1]

    target_variables = df[Class].unique()
    variables = df[attribute].unique()

    entropy2 = 0
    for variable in variables:
        entropy = 0
        for target_variable in target_variables:
            num = len(df[attribute][df[attribute] == variable][df[Class] == target_v
            den = len(df[attribute][df[attribute] == variable])
            fraction = num / (den+eps)
            entropy += -fraction * log(fraction+eps)
        fraction2 = den / len(df)
        entropy2 += -fraction2 * entropy

    return abs(entropy2)
```

In [16]:

```
# Calculate information gain and return the best splitting node (feature)
def infoGain(df):
    IG = []
    for key in df.keys()[:-1]:
        IG.append(findEntropy(df) - findEntropyAttribute(df, key))

    return df.keys()[:-1][np.argmax(IG)]
```

In [17]:

```
def giniImpurity2(valueCounts):
    n = valueCounts.sum()
    p_sum = 0
    for key in valueCounts.keys():
        p_sum = p_sum + (valueCounts[key] / n) * (valueCounts[key] / n)
        gini = 1 - p_sum

    return gini
```

In [18]:

```
# Calculating gini impurity for the attiributes
def giniSplitAtt2(df, attName):
    attValues = df[attName].value_counts()
    gini_A = 0
    for key in attValues.keys():
        dfKey = df[className][df[attName] == key].value_counts()
        numOfKey = attValues[key]
        n = df.shape[0]
        gini_A = gini_A + ((numOfKey / n) * giniImpurity2(dfKey))

    return gini_A
```

In [19]:

```
def getSubtable(df, node, value):
    return df[df[node] == value].reset_index(drop=True)
```

In [20]:

```
def giniIndex2(df, attributeNames):
    giniAttribute = {}
    minValue = sys.maxsize
    for key in attributeNames:
        giniAttribute[key] = giniSplitAtt2(df, key)
        if giniAttribute[key] < minValue:
            minValue = giniAttribute[key]
            selectedAttribute = key
    minValue = min(giniAttribute.values())
    return selectedAttribute
```

Tree Induction Function

In [21]:

```
def buildTree(df, model, tree=None):
    Class = df.keys()[-1]
    if model == 'infoGain':
        #print("building infoGain")
        node = infoGain(df)
    else:
        #print("building gini Index")
        node = giniIndex2(df, attName)
    attValueBT = np.unique(df[node])
    #Create an empty dictionary to create tree
    if tree is None:
        tree = {}
        tree[node] = {}
    #We make loop to construct a tree by calling this function recursively.
    #In this we check if the subset is pure and stops if it is pure.

    for value in attValueBT:
        #print('value (buildTree): ', value)
        subtable = getSubtable(df, node, value)

        clValue, counts = np.unique(subtable[className], return_counts=True)

        if len(counts) == 1:
            tree[node][value] = clValue[0]
        else:
            tree[node][value] = buildTree(subtable, model)
    return tree
```

Creating Decision Tree

Helper function to get training and testing split

In [12]:

```
def train_test_split(df, test_size):

    if isinstance(test_size, float):
        test_size = round(test_size * len(df))

    indices = df.index.tolist()
    test_indices = random.sample(population=indices, k=test_size)

    test_df = df.loc[test_indices]
    train_df = df.drop(test_indices)

    return train_df, test_df
```

In []:

```
train_df, test_df = train_test_split(pre_processed_df, test_size=0.4)#get 60% training
```

Splitting the data 60/40

In []:

```
test_df, valid_df = train_test_split(test_df, test_size=0.5)#get 20% 20% test and va
```

Splitting 40% of the data allocated to testing into half for a 20/20 test and validation split

In [91]:

```
attName = list(train_df.columns)[-1]
className = 'class'
```

Building the Information Gain and Gini Index Tree

In [93]:

```
info_gain_tree = buildTree(train_df, "infoGain")
```

In [95]:

```
gini_index_tree = buildTree(train_df, "giniIndex")
```

Accuracy of Information Gain Split Criteria Decision Tree Pre-Pruning

Helper Function for Decision Tree Classification and Accuracy Report

In [26]:

```
def accuracy_of_the_tree(instance, tree, default=None):
    attribute = list(tree.keys())[0]
    if instance[attribute] in tree[attribute].keys():
        result = tree[attribute][instance[attribute]]
        if isinstance(result, dict):
            return accuracy_of_the_tree(instance, result)
        else:
            return result
    else:
        return default
```

In [27]:

```
def _unique(seq, return_counts=False, id=None):

    found = set()
    if id is None:
        for x in seq:
            found.add(x)

    else:
        for x in seq:
            x = id(x)
            if x not in found:
                found.add(x)
    found = list(found)
    counts = [seq.count(0), seq.count(1)]
    if return_counts:
        return found, counts
    else:
        return found

def _sum(data):
    sum = 0
    for i in data:
        sum = sum + i
    return sum
```

Accuracy of Information Gain Split Criteria Decision Tree Pre Pruning

Training set

In [126]:

```
temptrain1 = pd.DataFrame()
temptrain1['predicted'] = train_df.apply(accuracy_of_the_tree, axis=1, args=(info_ga
print( 'Accuracy with info gain ' + (str( sum(train_df['class']==temptrain1['predic
```

Accuracy with info gain 63.55448814169919

Testing set

In [127]:

```
temptest1 = pd.DataFrame()
temptest1['predicted'] = test_df.apply(accuracy_of_the_tree, axis=1, args=(info_gain,))
print( 'Accuracy with info gain ' + (str( sum(test_df['class']==temptest1['predicted'])))
```

Accuracy with info gain 63.88570492877027

Findings

We can observe that the accuracy of the training set and testing set are hovering around 63% accuracy, therefore we can conclude the no overfitting has occurred and the model has achieved a decent level of generalisation.

Accuracy of Gini Index Split Criteria Decision Tree Pre Pruning

Training set

In [128]:

```
temptrain1 = pd.DataFrame()
temptrain1['predicted2'] = train_df.apply(accuracy_of_the_tree, axis=1, args=(gini_index,))
print( 'Accuracy with gini index ' + (str( sum(train_df['class']==temptrain1['predicted2'])))
```

Accuracy with gini index 63.63090527005267

Testing set

In [129]:

```
temptest1 = pd.DataFrame()
temptest1['predicted2'] = test_df.apply(accuracy_of_the_tree, axis=1, args=(gini_index,))
print( 'Accuracy with gini index ' + (str( sum(test_df['class']==temptest1['predicted2'])))
```

Accuracy with gini index 63.91026690682823

Findings

We can observe that the accuracy of both the training and testing data set are hover around 63% as well. Therefore, the tree using the gini index split criteria does not have an overfitting issues and has obtained generalisation

Final conclusion of both models

Both split criteria of Information Gain and Gini Index Yielded similar accuracies of 63% with no signs of overfitting and displayed the ability to generalise for unseen data. We will therefore attempt to tune and improve its performance but Post Pruning the decision tree in hopes of reducing computational complexity and

model accuracy.

Post Pruning Helper Funtion

In [199]:

```
def preorder (temptree, number):
    if isinstance(temptree, dict):
        attribute = list(temptree.keys())[0]
        if temptree[attribute]['number'] == number:
            if(temptree[attribute][0]!=0 and temptree[attribute][0]!=1):
                temp_tree = temptree[attribute][0]
                if isinstance(temp_tree, dict):
                    temp_attribute = list(temp_tree.keys())[0]
                    temptree[attribute][0] = temp_tree[temp_attribute]['best_class']
            elif(temptree[attribute][1]!=0 and temptree[attribute][1]!=1):
                temp_tree = temptree[attribute][1]
                if isinstance(temp_tree, dict):
                    temp_attribute = list(temp_tree.keys())[0]
                    temptree[attribute][1] = temp_tree[temp_attribute]['best_class']
        else:
            left = temptree[attribute][0]
            right = temptree[attribute][1]
            preorder(left, number)
            preorder(right, number)
    return temptree
```

In [254]:

```
def count_number_of_non_leaf_nodes(tree):
    if isinstance(tree, dict):
        attribute = list(tree.keys())[0]
        left = tree[attribute][0]
        print(left)
        right = tree[attribute][1]
        print(right)
        return (1 + count_number_of_non_leaf_nodes(left) + count_number_of_non_leaf_nodes(right))
    else:
        return 0;
```

In [201]:

```

def post_prune(L, K, tree):
    best_tree = tree
    for i in range(1, L+1) :
        temp_tree = copy.deepcopy(best_tree)
        M = randint(1, K);
        for j in range(1, M+1):
            n = count_number_of_non_leaf_nodes(temp_tree)
            if n > 0:
                P = randint(1,n)
            else:
                P = 0
            preorder(temp_tree, P)
        test_data['accuracyBeforePruning'] = test_data.apply(accuracy_of_the_tree, a
        accuracyBeforePruning = str( sum(test_data['Class']==test_data['accuracyBefo
        test_data['accuracy_after_pruning'] = test_data.apply(accuracy_of_the_tree,
        accuracy_after_pruning = str( sum(test_data['Class']==test_data['accuracy_af
        if accuracy_after_pruning >= accuracyBeforePruning:
            best_tree = temp_tree
    return best_tree

```

Accuracy Report for post pruning

I am unable to proceed as the functions created were generating and error due to the multi way split nature of my tree induction function