CSCI316 Software and Programming Packages Installation Guide to install and Run PySpark in Jupyter Notebook on Windows (Adopted from instructions provided by Dr Guoxin Su)

(SJ: 25 June 2021)

Installing and Integrating PySpark with Jupyter Notebook on Windows

A. Downloading PySpark, Anaconda, 7-Zip, and Java JDK.

Before setting up PySpark, we need to have the following packages installed in the System:

- **Download PySpark package**
  i. https://spark.apache.org/downloads.html

  ii. For the package type, please choose **'Pre-built for Apache Hadoop 2.7'**
  iii. Click the label **'spark-3.1.2-bin-hadoop2.7.tgz'** to start downloading the package.
  iv. In the next screen, choose the mirror site that you want to download the package from, and double-click on the link to start downloading.

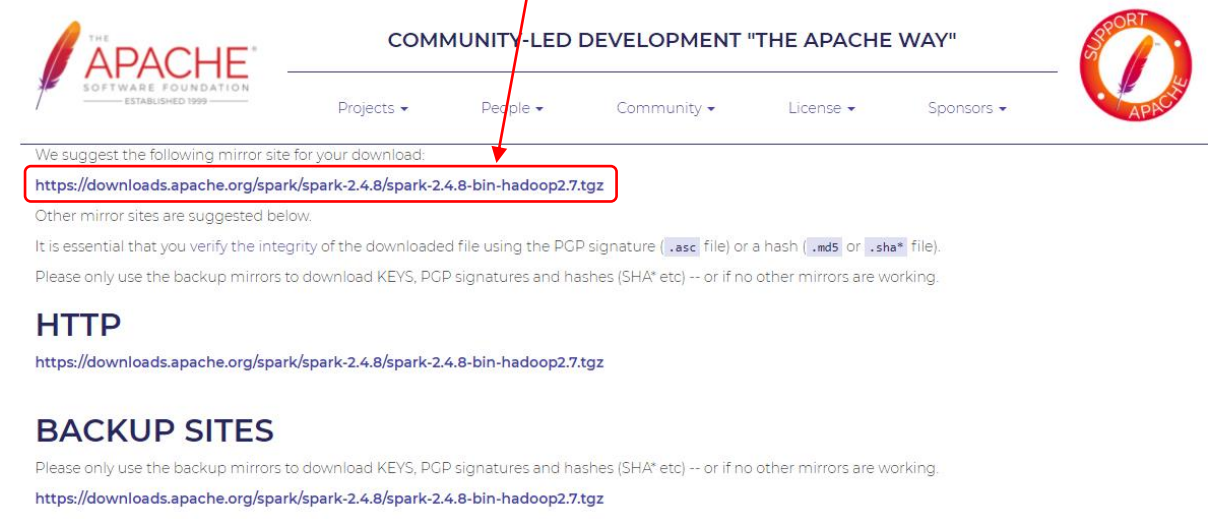

  v. Save the downloaded package to a working directory.

- **Installing 7-zip**
  i. If you already have 7-zip installed on your Windows, you can skip this step, otherwise, download the 7-zip installer and install the 7-zip application.
  ii. https://www.7-zip.org/download.html

- Choose the version of installer that fits your system and click the 'Download' link to start downloading the installer. For example, for my system, I choose the **64-bit x64** version.

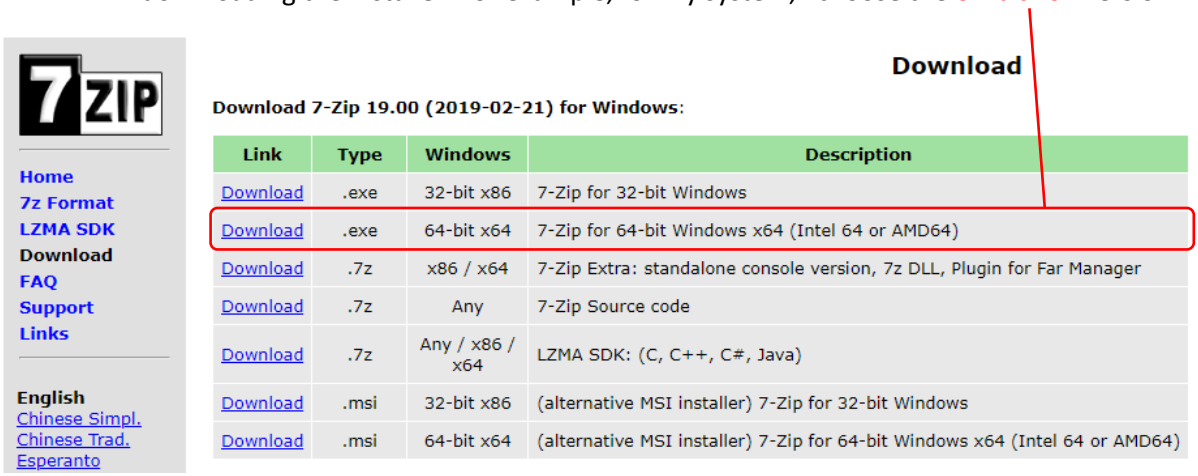

7-Zip

Home
7z Format
LZMA SDK
Download
FAQ
Support
Links

English
Chinese Simpl.
Chinese Trad.
Esperanto

Download

Download 7-Zip 19.00 (2019-02-21) for Windows:

Link	Type	Windows	Description
Download	.exe	32-bit x86	7-Zip for 32-bit Windows
Download	.exe	64-bit x64	7-Zip for 64-bit Windows x64 (Intel 64 or AMD64)
Download	.7z	x86 / x64	7-Zip Extra: standalone console version, 7z DLL, Plugin for Far Manager
Download	.7z	Any	7-Zip Source code
Download	.7z	Any / x86 / x64	LZMA SDK: (C, C++, C#, Java)
Download	.msi	32-bit x86	(alternative MSI installer) 7-Zip for 32-bit Windows
Download	.msi	64-bit x64	(alternative MSI installer) 7-Zip for 64-bit Windows x64 (Intel 64 or AMD64)

- Once the downloading is completed, and the installer is checked, execute the installer to install the 7-Zip application. You can follow all the default setting unless you want to install the 7-Zip differently.

- **Installing Java JDK**
  i. If your system has no Java installed or the Java version is 7.x or less, please download and install Java from Oracle: https://www.oracle.com/java/technologies/javase/javase-jdk8-downloads.html
  ii. For my system, I downloaded and install Java SE Development Kit (Java JDK 8) (Note: So far this version works, I have tried installing the latest JDK and it did not work.)

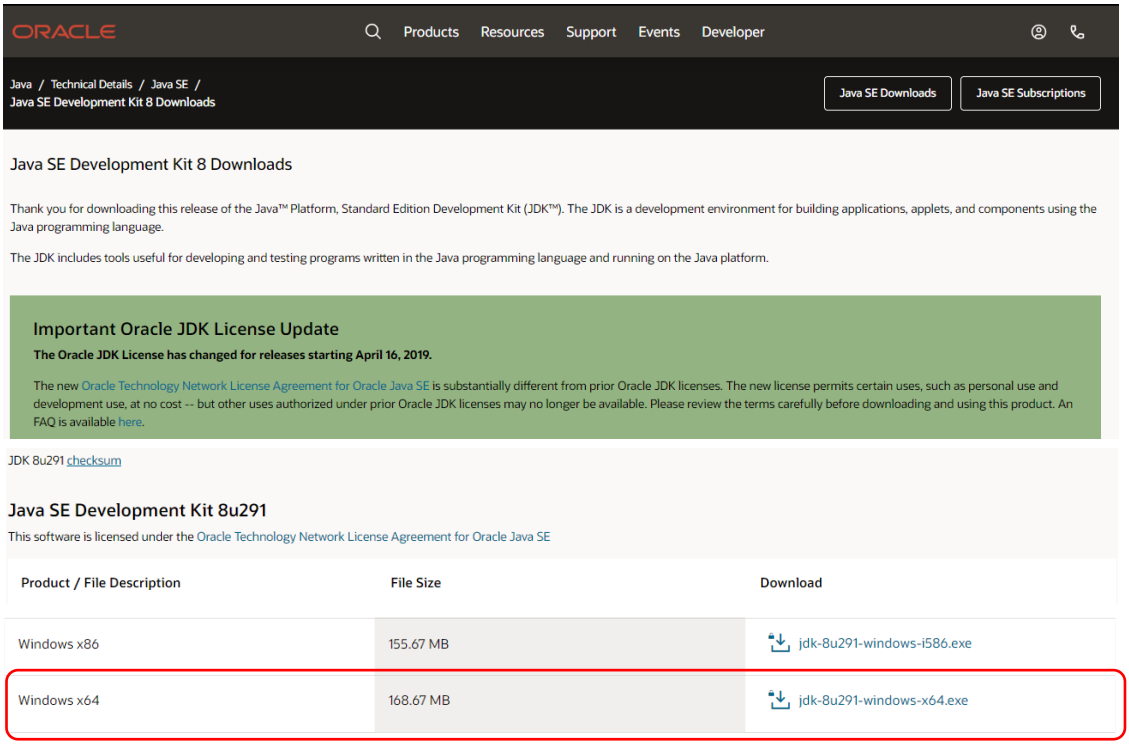

ORACLE Products Resources Support Events Developer

Java / Technical Details / Java SE / Java SE Development Kit 8 Downloads

Java SE Downloads | Java SE Subscriptions

Java SE Development Kit 8 Downloads

Thank you for downloading this release of the Java™ Platform, Standard Edition Development Kit (JDK™). The JDK is a development environment for building applications, applets, and components using the Java programming language.

The JDK includes tools useful for developing and testing programs written in the Java programming language and running on the Java platform.

Important Oracle JDK License Update

The Oracle JDK License has changed for releases starting April 16, 2019.

The new Oracle Technology Network License Agreement for Oracle Java SE is substantially different from prior Oracle JDK licenses. The new license permits certain uses, such as personal use and development use, at no cost -- but other uses authorized under prior Oracle JDK licenses may no longer be available. Please review the terms carefully before downloading and using this product. An FAQ is available here.

JDK 8u291 checksum

Java SE Development Kit 8u291

This software is licensed under the Oracle Technology Network License Agreement for Oracle Java SE

Product / File Description	File Size	Download
Windows x86	155.67 MB	jdk-8u291-windows-i586.exe
Windows x64	168.67 MB	jdk-8u291-windows-x64.exe

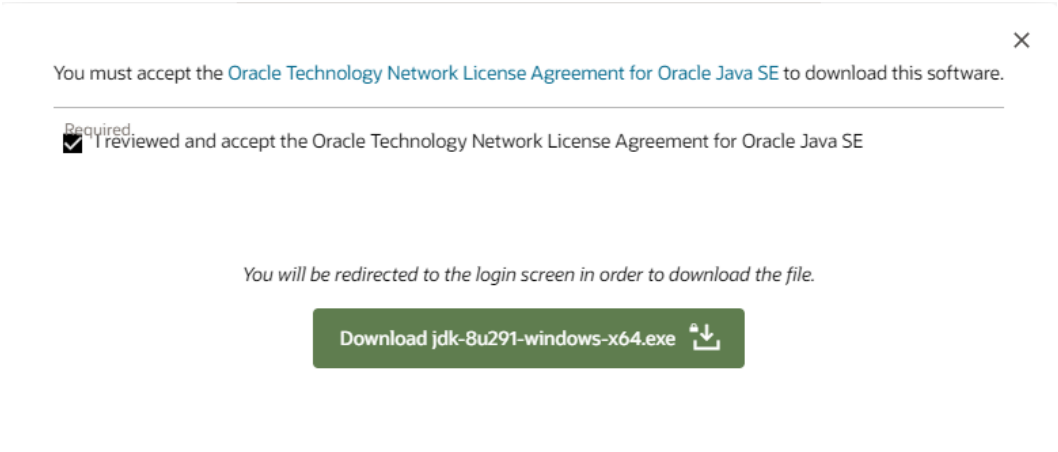

- Save to a directory of your choice.

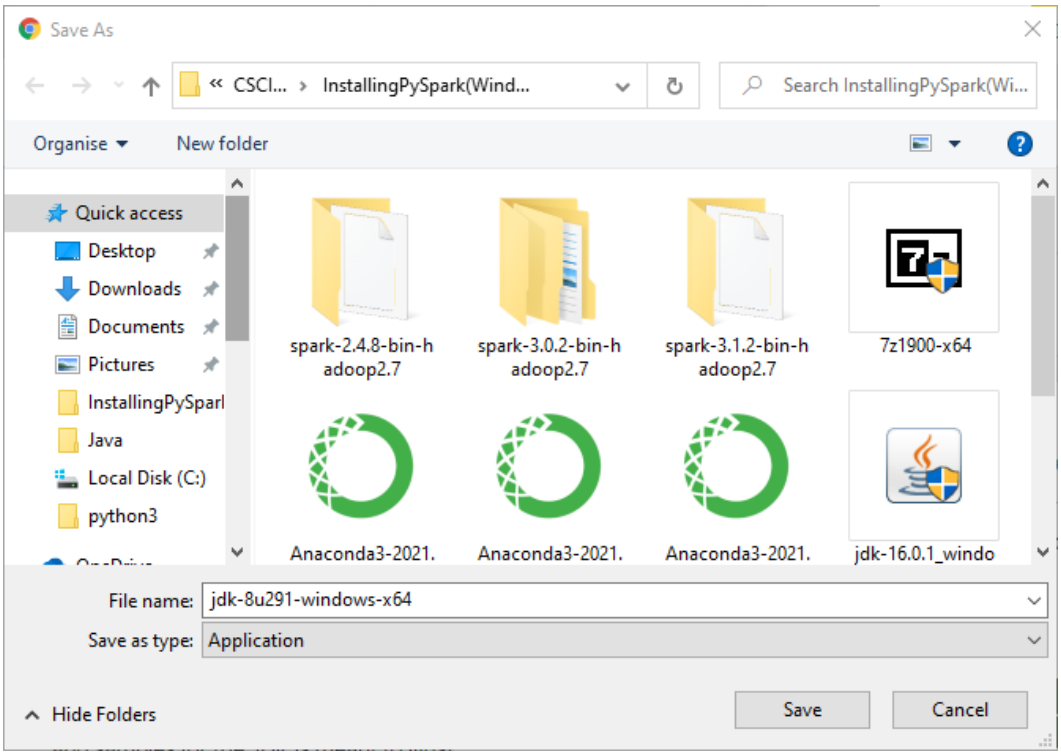

- Next, you can proceed to install the Java JDK, but please do not use the default setting from the installer.

- Installing Java JDK:
    i. Create a new folder in your local drive; it can be in ‘C:\’ or ‘D:\’. For me, I install it in my ‘C:\’ drive, and name the folder ‘Java’. Note: Please use a name that does not have a space in between.

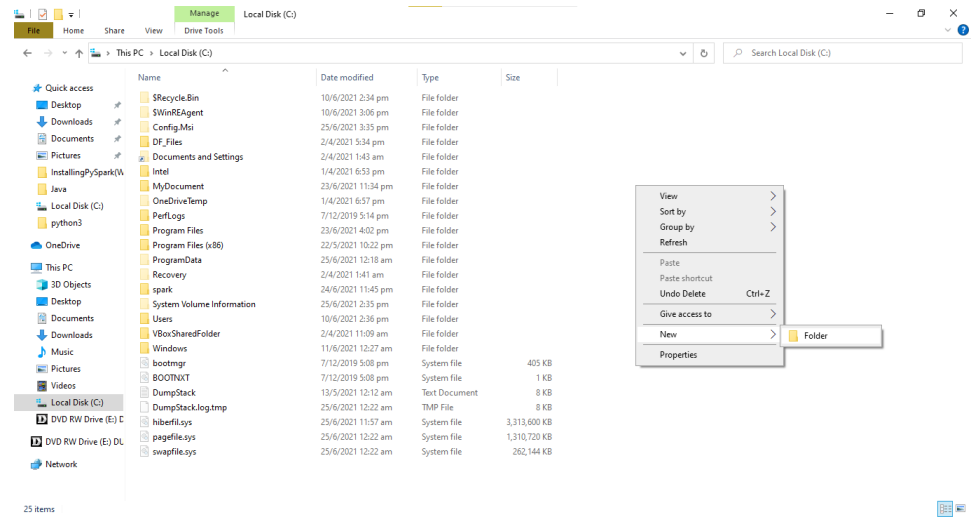

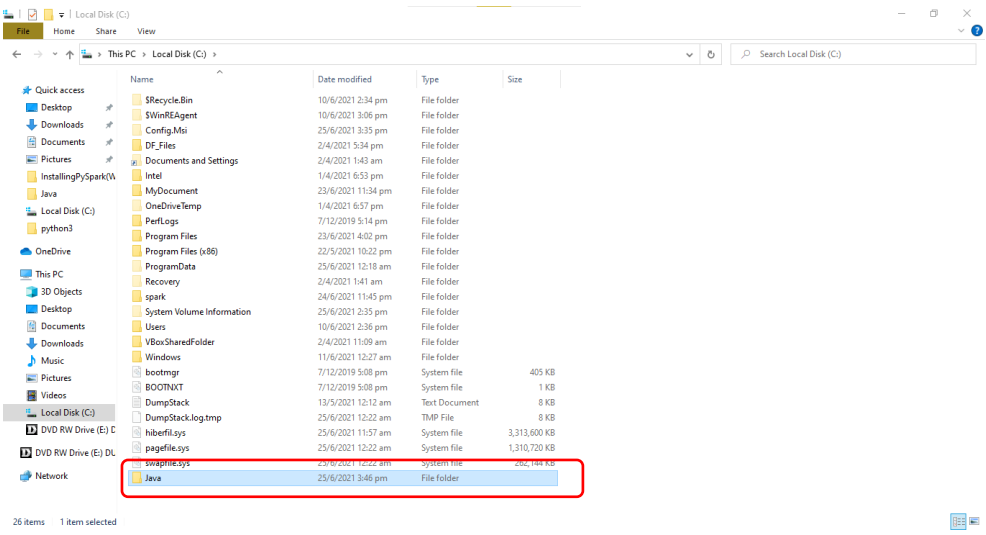

ii. Execute your java JDK installer. From the directory where you save the Java JDK installer, **double-click** on the installer to start the installation.

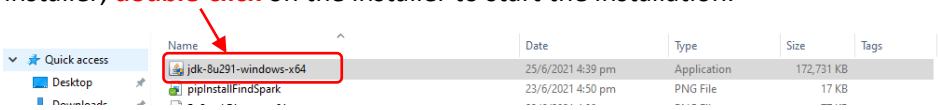

iii. In the pop-up menu, click 'Next' to proceed.

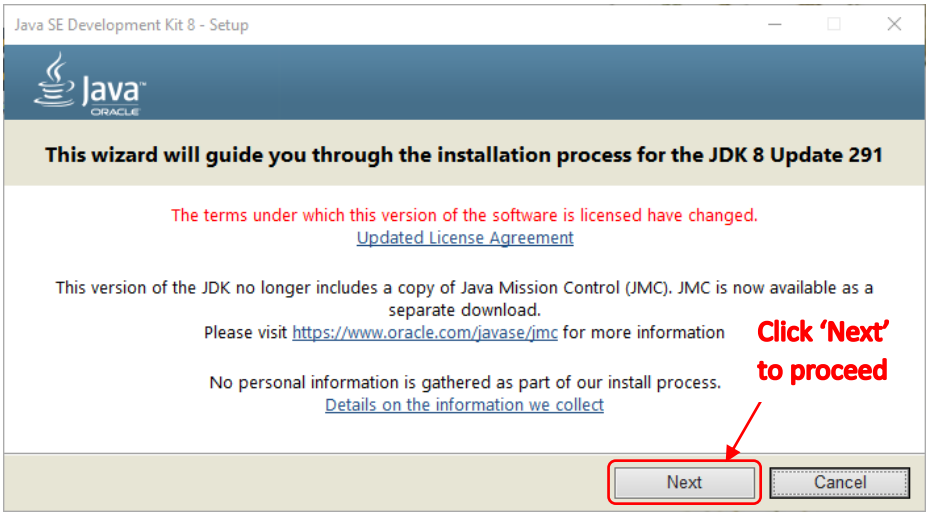

iv. Click the 'Change…' button to change the directory where you intend to install your Java JDK.

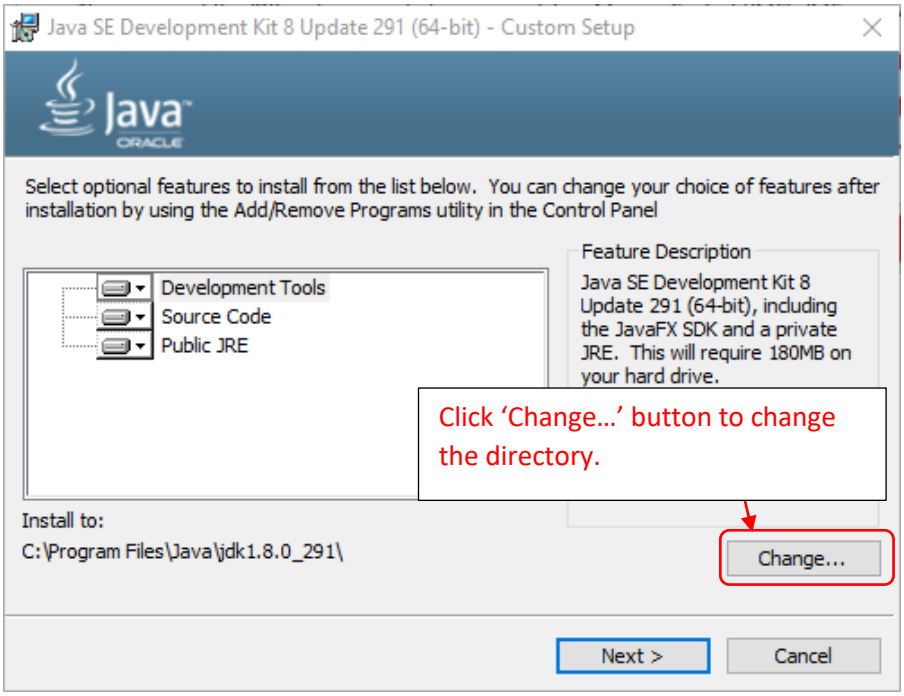

v. In the next window, click on the icon to proceed to change the directory.

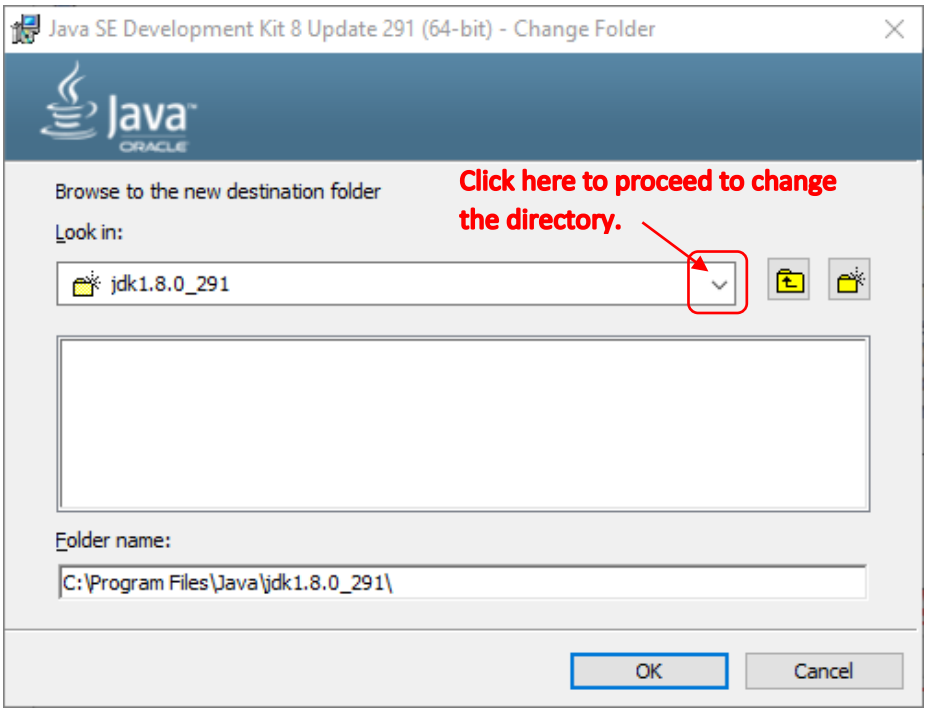

vi. Change to the directory where you want to install you Java JDK and click the 'OK' button.

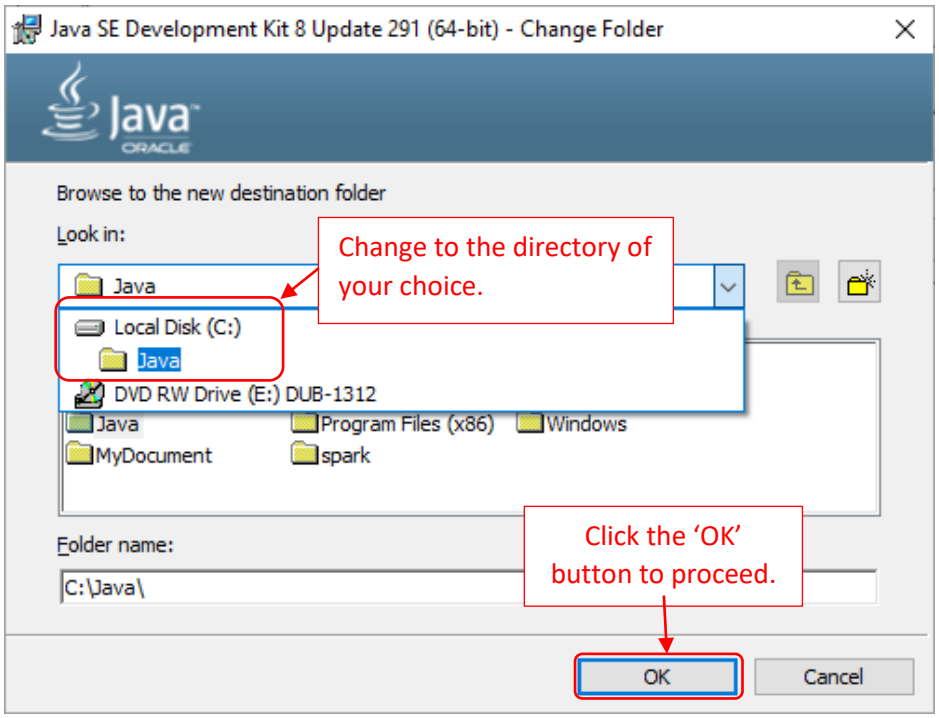

vii. Click on the ‘Next’ button to continue the installation.

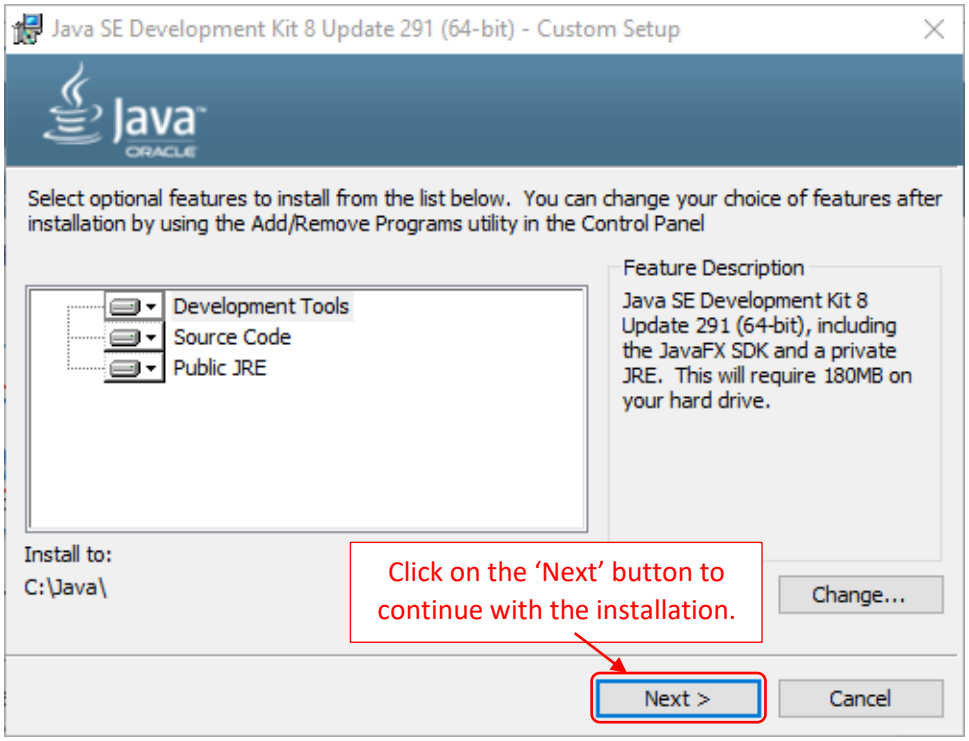

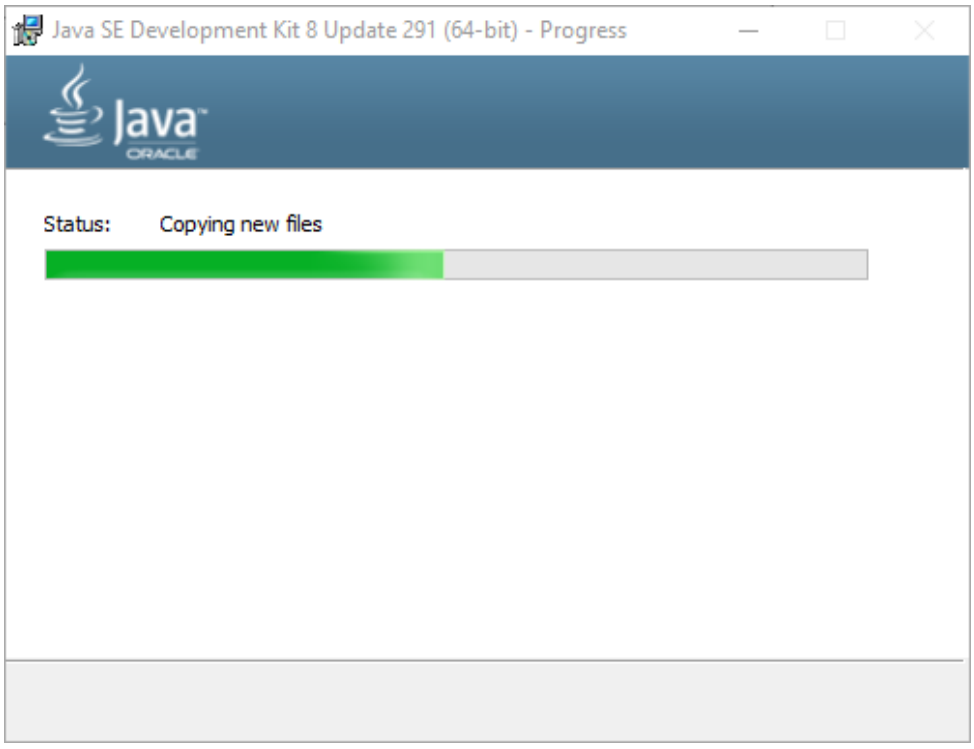

viii. In the next window, similarly, click the 'Change…' button to change to a different directory to install the Java jre.

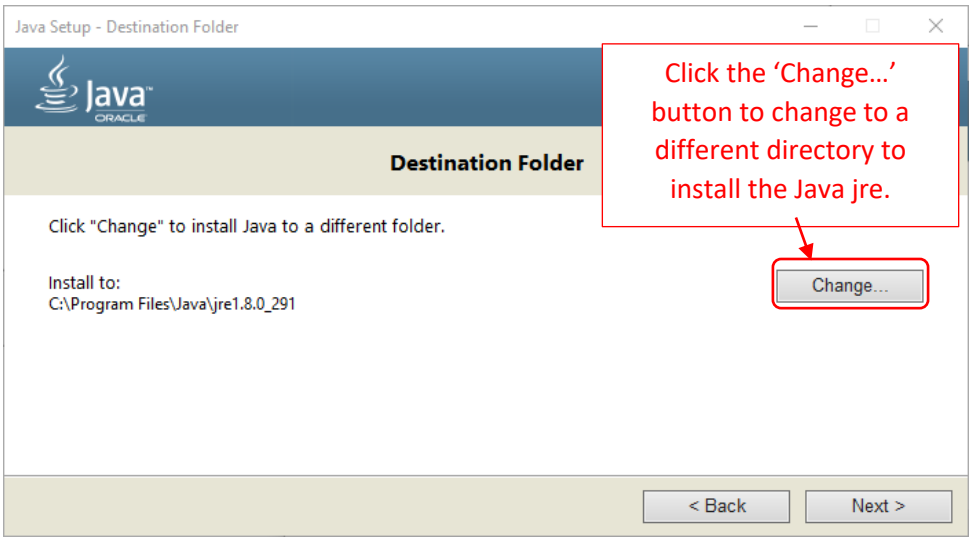

ix. Change to the directory where you want to install you Java JRE. You may want to create (make) a new directory for that.

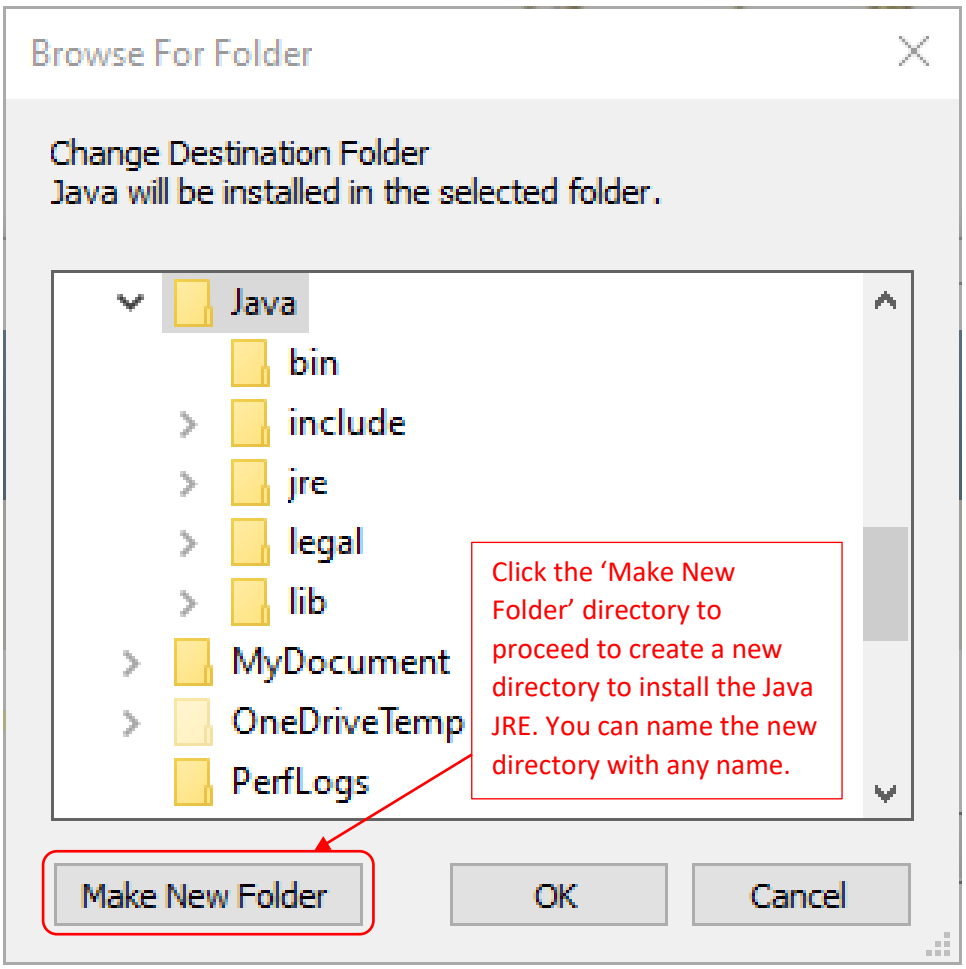

x. Click on the newly created directory (in my example, JavaJRE180_291) and click the ‘OK’ button to continue.

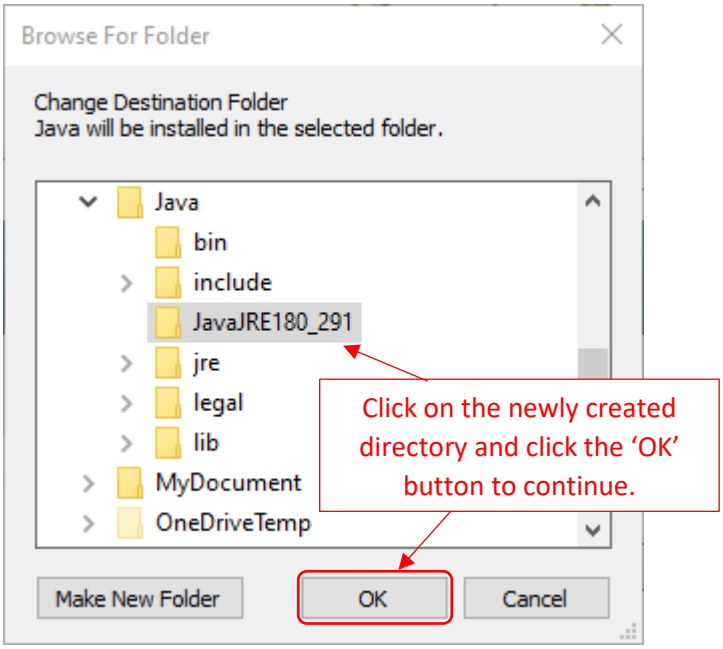

xi. Check that the directory is correct and click the ‘Next’ button to proceed with the installation.

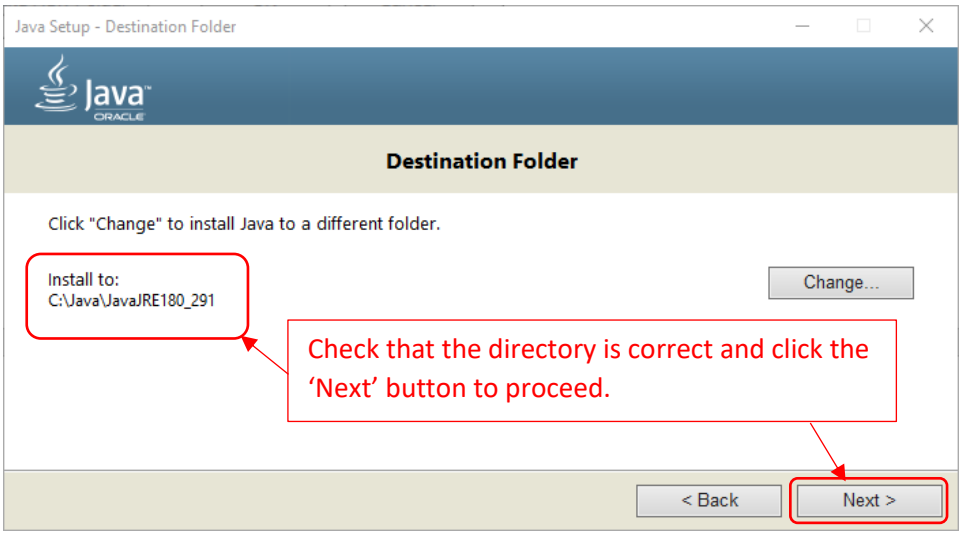

xii. Wait for the installation to complete….

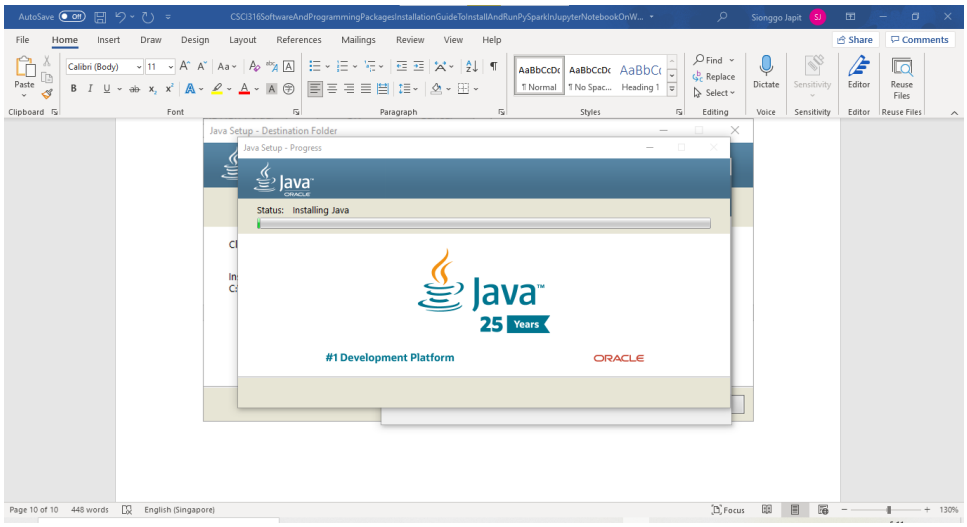

xiii. Click the ‘Close’ button when done.

- **Installing Python and Jupyter Notebook**
  i. You can get both Python and Jupyter by installing the Python 3.x version of Anaconda distribution.
  ii. https://www.anaconda.com/products/individual
  iii. Scroll all the way down to the bottom of the page and choose the version of the installer that suit your system to start the download.

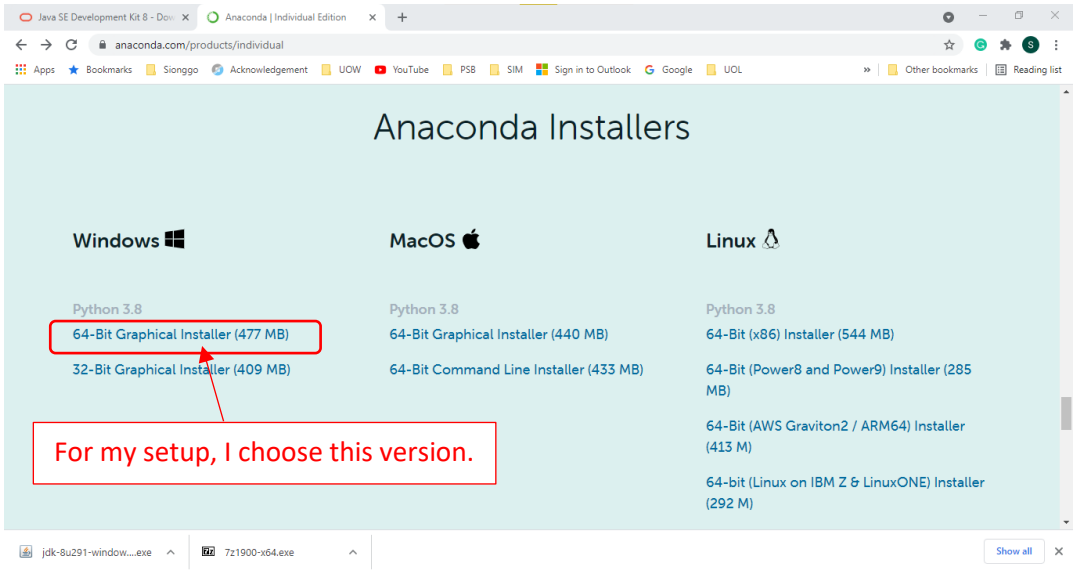

  iv. Click the 'Download' button to start the download.
  v. Installing Anaconda:
    - After the downloading is complete, you can proceed to install the Anaconda distribution.

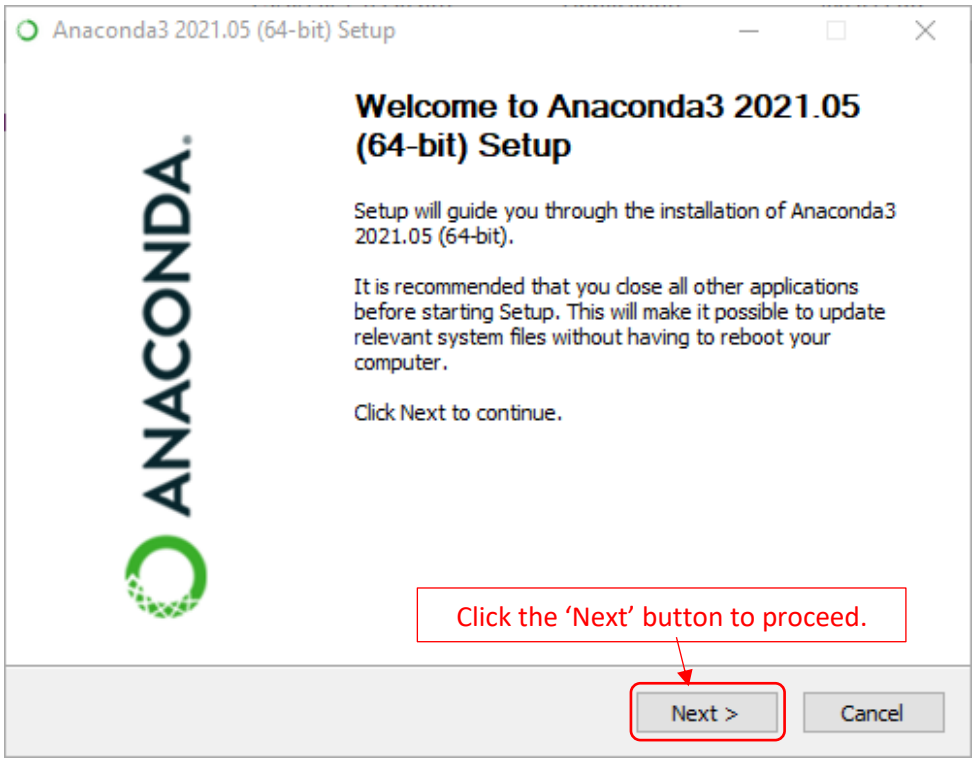

- Read the ‘License Agreement’ and when finished reading, click the ‘I Agree’ button to proceed with the installation if you agree with the terms and condition specified in the License Agreement. 😊.

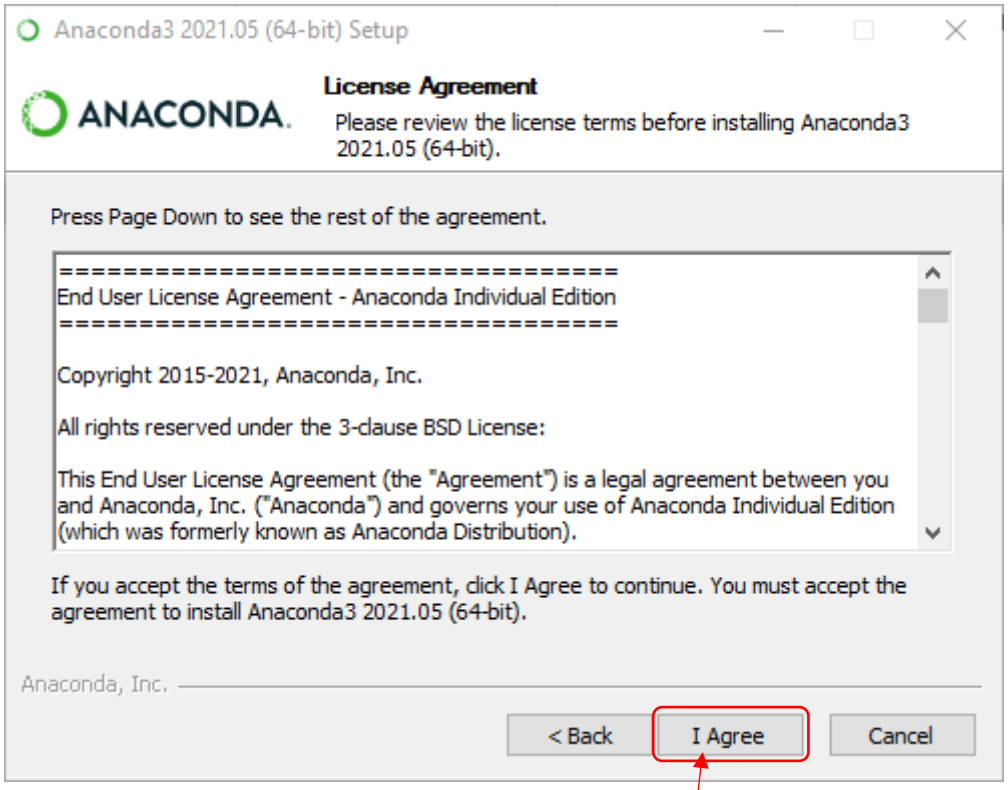

Click the ‘I Agree’ button to proceed with the installation if you agree with the terms and condition specified in the License Agreement.

- Select the radio button ‘Just Me’ and click the ‘Next’ button to proceed.

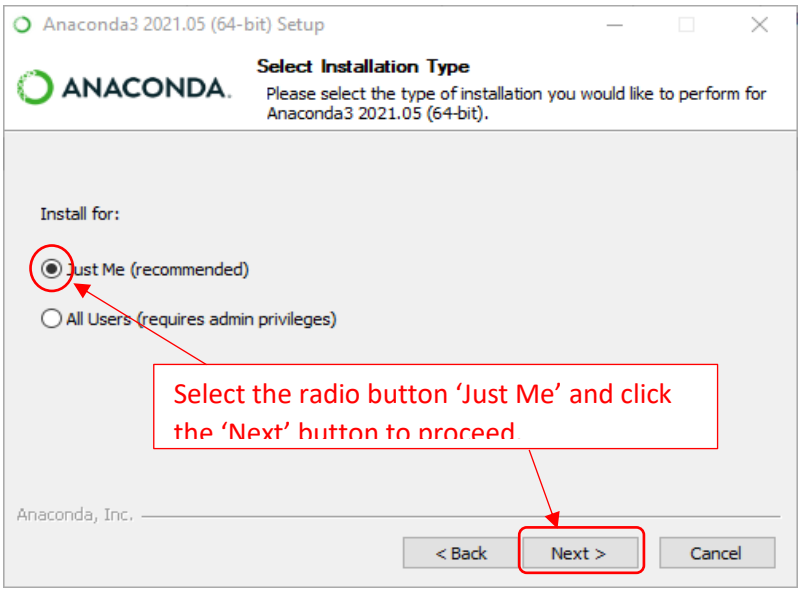

- You can accept the proposed destination folder. You can also change to a different directory if you want to. Click the 'Next' button to proceed.

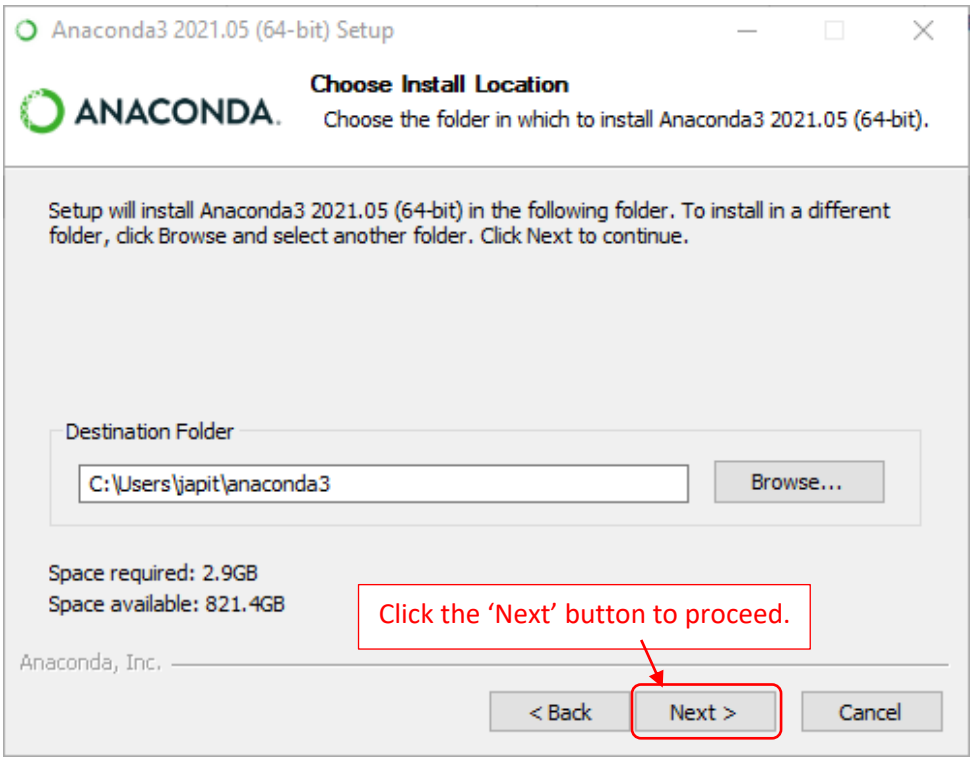

- Check the 'Register Anaconda3 as my default Python 3.8' option and click the 'Install' button to start the installation. This installation will take a while. Just relax….

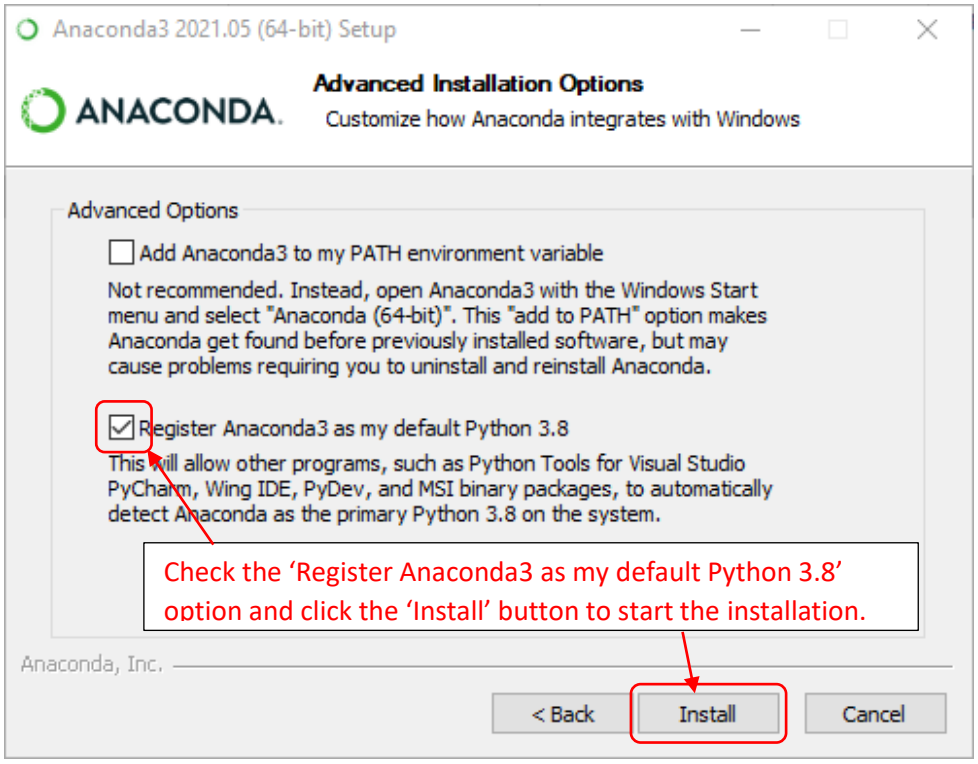

B. Installing PySpark

- While waiting for the Anaconda to finish the installation, we can proceed to un-zip the PySpark package for the PySpark setup.
    i. Unpack the spark package 'spark-3.1.2-bin-hadoop2.7.tgz' that was downloaded earlier. I save the package in my drive C:\Users\japit\Teaching\UOW\CSCI316\InstallingPySpark(Windows)\ spark-3.1.2-bin-hadoop2.7.tgz.

    This package is rather special. It has been packed (zipped) two times. Hence, you need to unpack (unzip) the package two times.

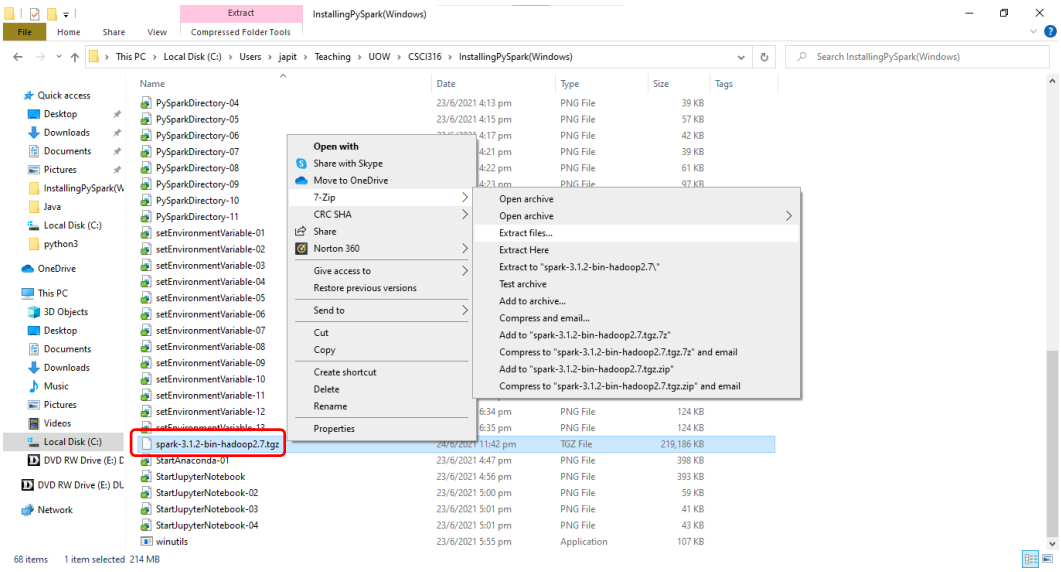

Click the 'OK' button to proceed.

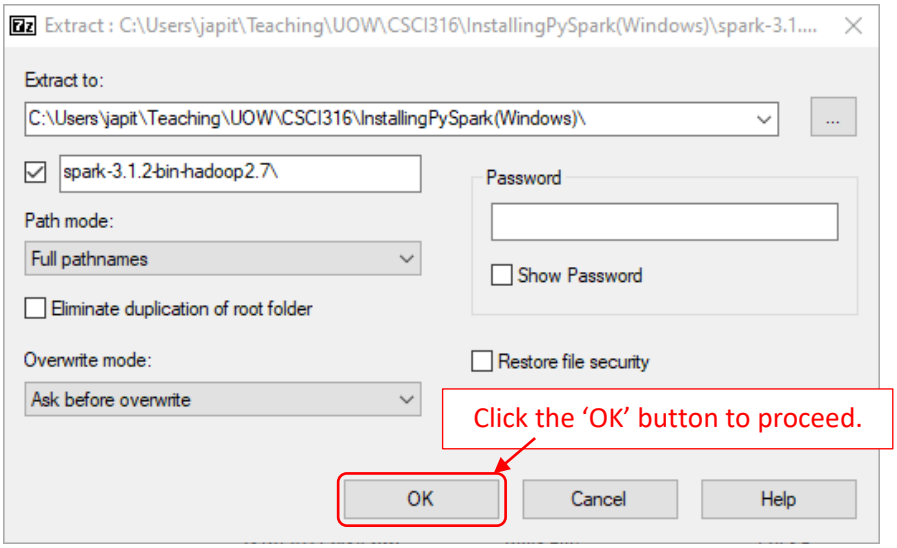

After the unpacking, navigate to the folder of the unpacked file, and do another unpacking process.

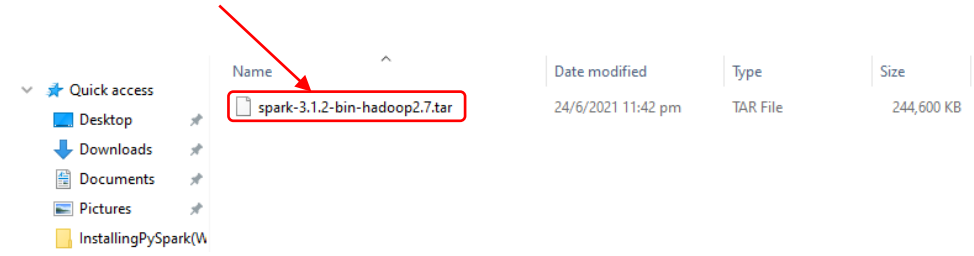

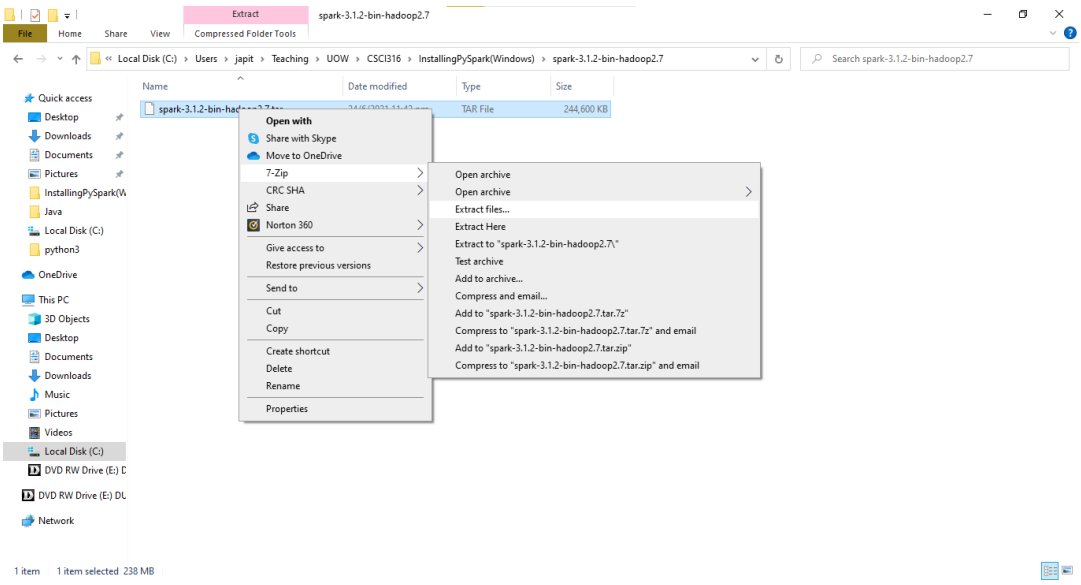

Click the 'OK' button to proceed with the unpacking process.

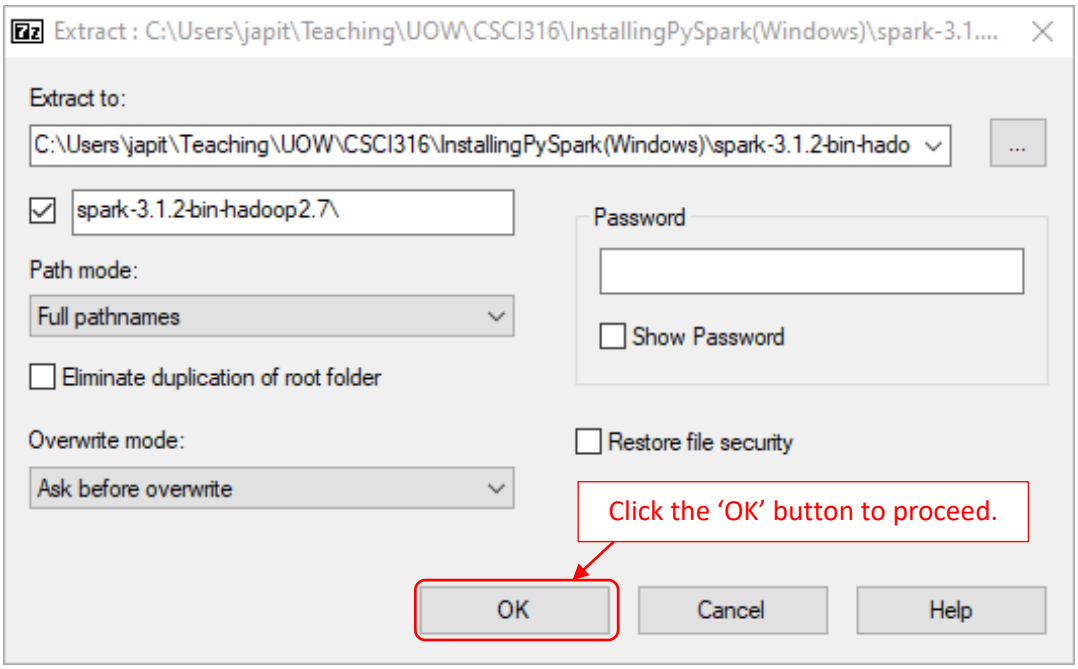

ii. Put the unpacked package to a directory of your choice. I put mine under C:\spark\spark-3.1.2-bin-hadoop2.7.

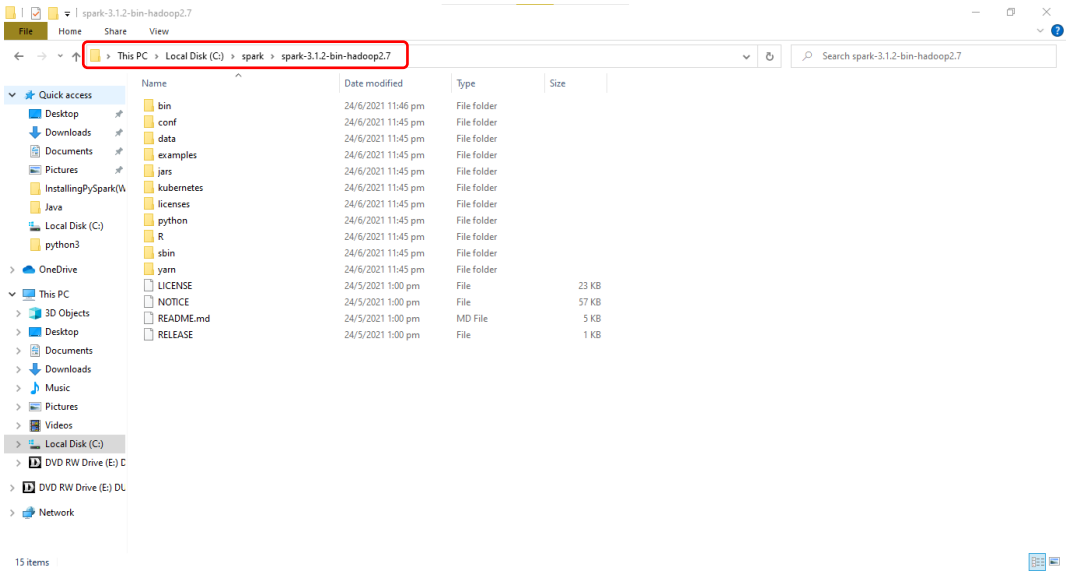

iii. 'winutils' is a collection of useful TCL commands that access some part of the Win32 API. This enables the user to use Windows specific services. Winutils is required when installing Hadoop on Windows environment. Winutils can be downloaded from Steve Loughran's GitHub repo.

- https://github.com/steveloughran/winutils/
- Go to the corresponding Hadoop version in the Spark distribution and find winutils.exe under /bin. Note that when we download the PySpark package in the earlier step, we choose the 'Pre-built for Apache Hadoop 2.7'.
- Click the link 'hadoop-2.7.1' to navigate to the next page.

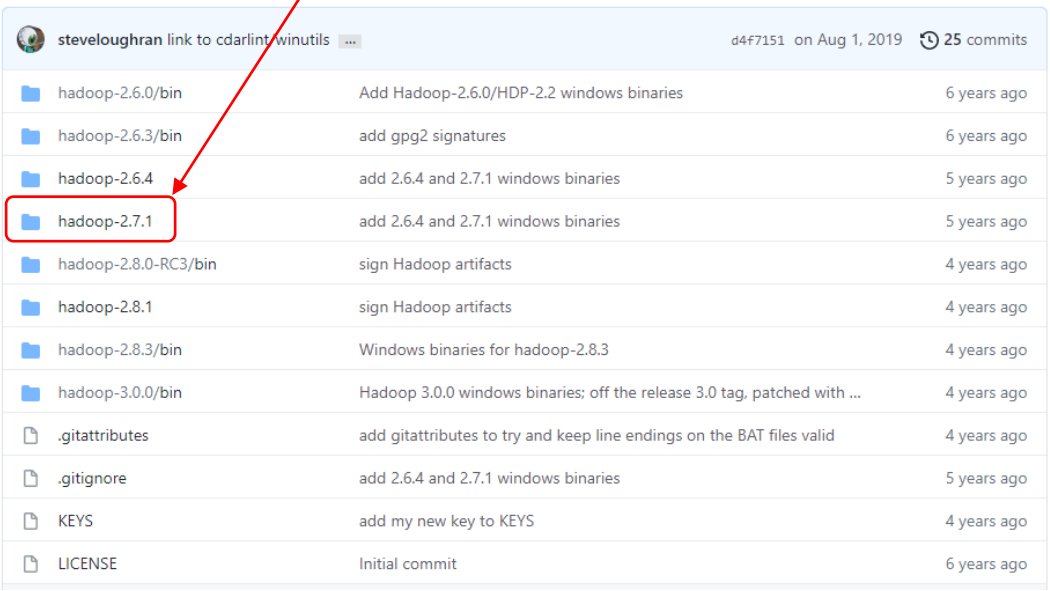

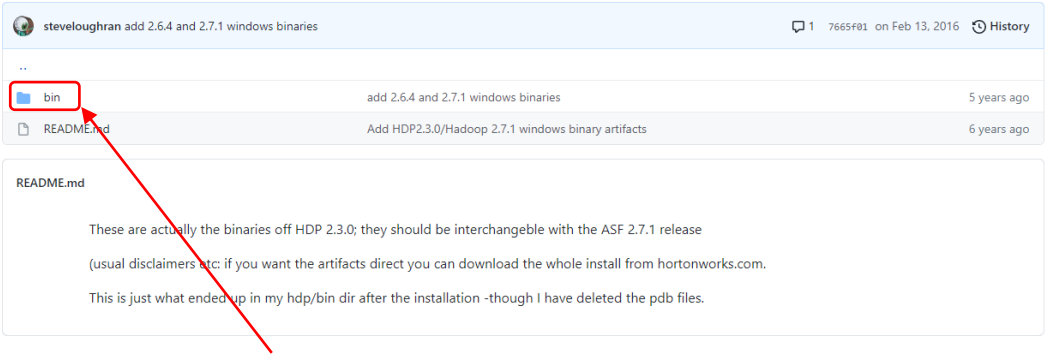

- Click the link ‘bin’ to navigate to the next page.

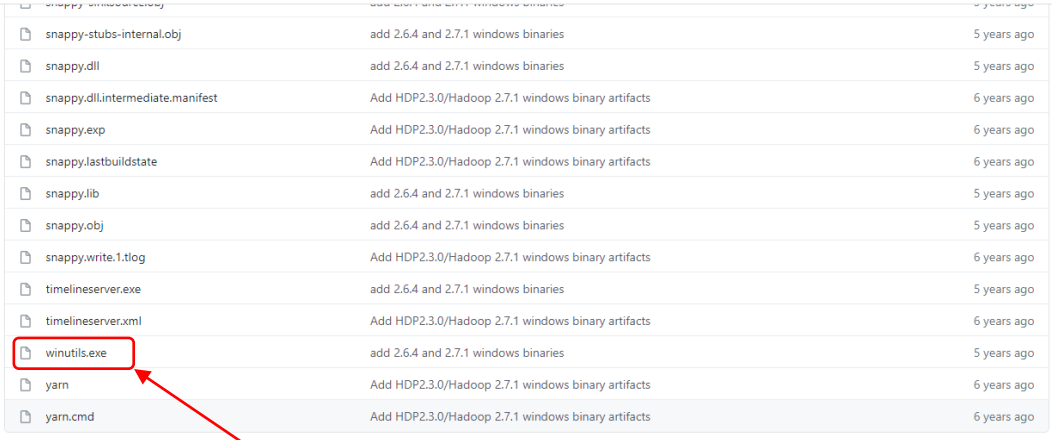

- Click the link ‘winutils.exe’ to navigate to the next page.

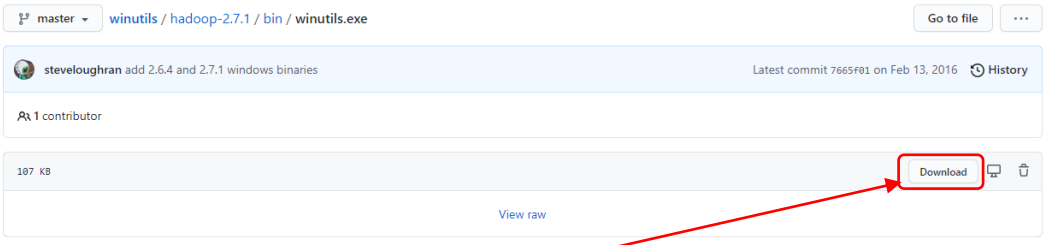

- Click the button ‘Download’ to start the download.
- The downloaded ‘winutils.exe’ need to be placed in the bin directory of the spark folders ‘C:\spark\spark-3.1.2-bin-hadoop2.7.’.

- The ‘winutils.exe’ is placed here.

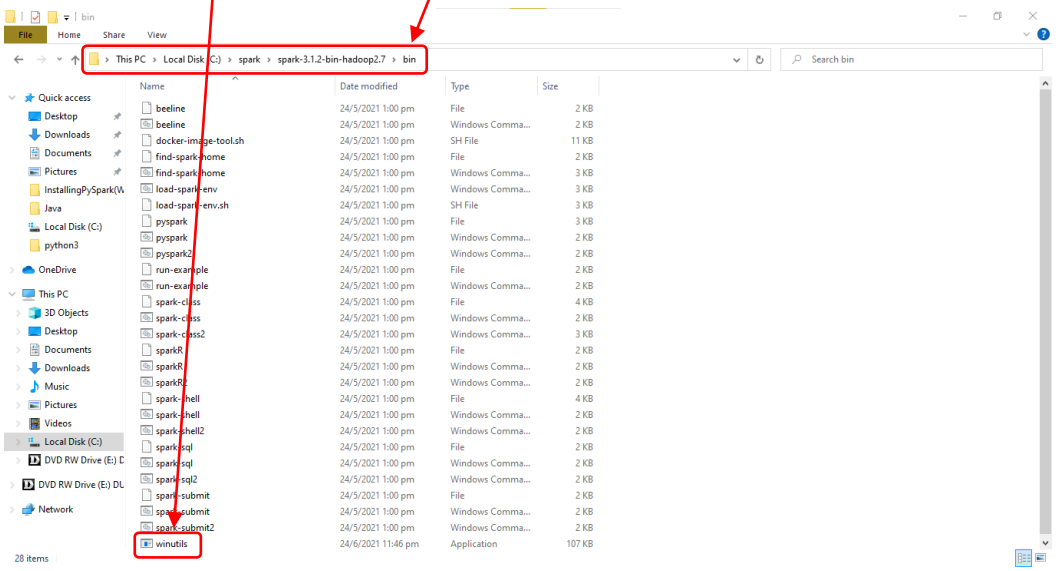

- **Setting of user’s environment variable**
- In order for the system to know where the various components (packages) are installed, we need to add the path (directory) of each component to the user’s environment variables. These variables let Windows find where the files are when we start the PySpark kernel.
- You can find the environment variable settings by typing “environment…” in the search box. When the System Properties windows is shown, click on the ‘Environment Variables…’ button to proceed.

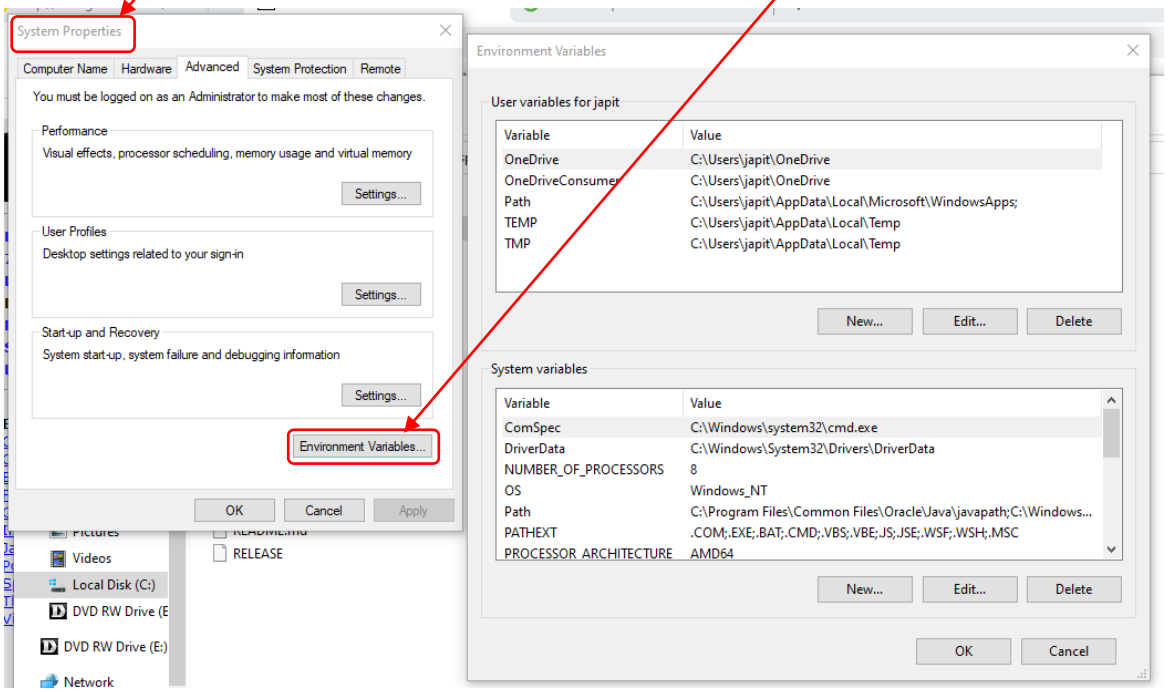

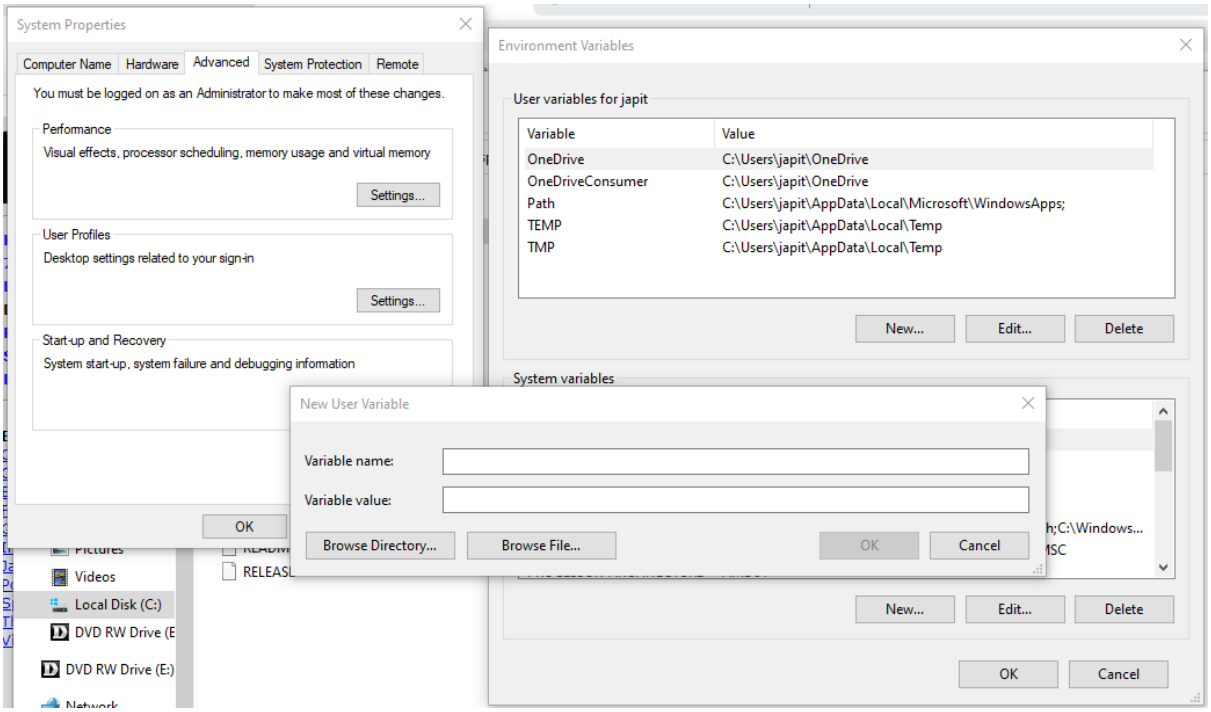
System Properties
Computer Name
Hardware
Advanced
System Protection
Remote
You must be logged on as an Administrator to make most of these changes.
Performance
Visual effects, processor scheduling, memory usage and virtual memory
Settings...
User Profiles
Desktop settings related to your sign-in
Settings...
Start-up and Recovery
System start-up, system failure and debugging information
OK
Environment Variables
User variables for japit
Variable
Value
OneDrive
C:\Users\japit\OneDrive
OneDriveConsumer
C:\Users\japit\OneDrive
Path
C:\Users\japit\AppData\Local\Microsoft\WindowsApps;
TEMP
C:\Users\japit\AppData\Local\Temp
TMP
C:\Users\japit\AppData\Local\Temp
New...
Edit...
Delete
System variables
New...
Edit...
Delete
OK
Cancel
New User Variable
Variable name:
Variable value:
Browse Directory...
Browse File...
OK
Cancel
Videos
Local Disk (C:)
DVD RW Drive (E:)
DVD RW Drive (E:)

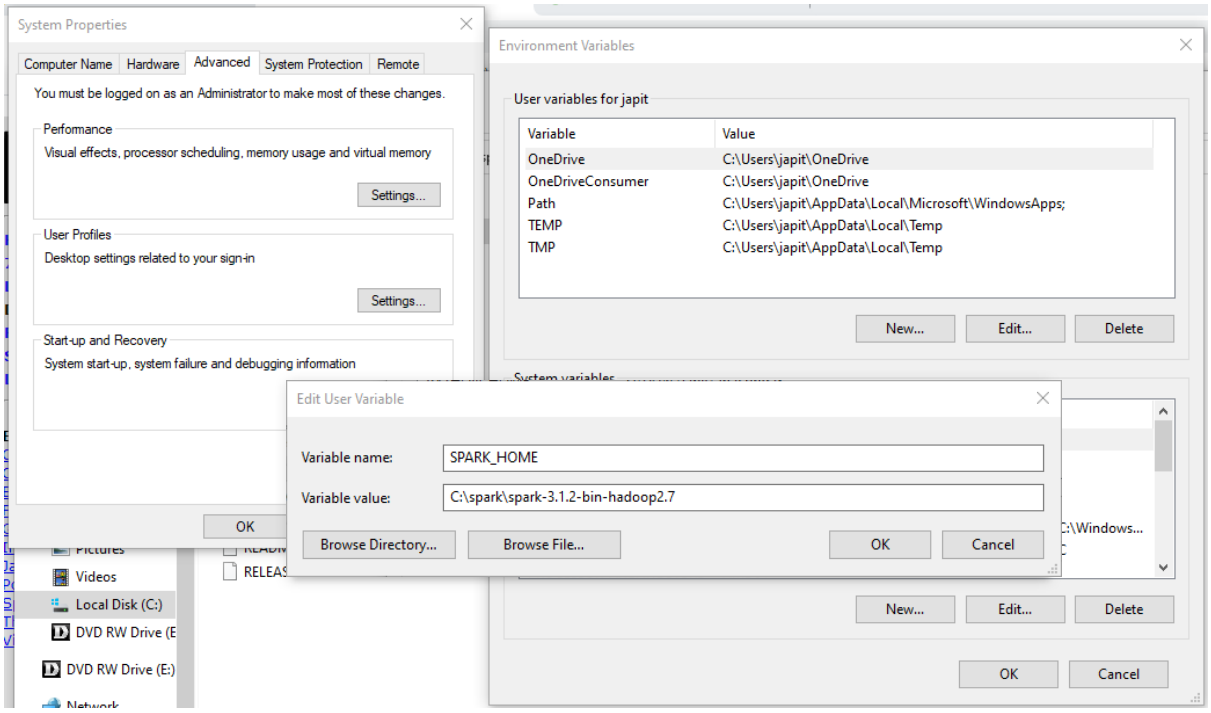
System Properties
Computer Name
Hardware
Advanced
System Protection
Remote
You must be logged on as an Administrator to make most of these changes.
Performance
Visual effects, processor scheduling, memory usage and virtual memory
Settings...
User Profiles
Desktop settings related to your sign-in
Settings...
Start-up and Recovery
System start-up, system failure and debugging information
OK
Environment Variables
User variables for japit
Variable
Value
OneDrive
C:\Users\japit\OneDrive
OneDriveConsumer
C:\Users\japit\OneDrive
Path
C:\Users\japit\AppData\Local\Microsoft\WindowsApps;
TEMP
C:\Users\japit\AppData\Local\Temp
TMP
C:\Users\japit\AppData\Local\Temp
New...
Edit...
Delete
New...
Edit...
Delete
OK
Cancel
Edit User Variable
Variable name:
SPARK_HOME
Variable value:
C:\spark\spark-3.1.2-bin-hadoop2.7
Browse Directory...
Browse File...
OK
Cancel
Videos
Local Disk (C:)
DVD RW Drive (E:)
DVD RW Drive (E:)

- Do the same for the rest of the components shown below. Note that the values shown here are according to my setup (the directories where I place the packages.) You need to follow according to your setup.

Name	Values
SPARK_HOME	C:\spark\spark-3.1.2-bin-hadoop2.7
HADOOP_HOME	C:\spark\spark-3.1.2-bin-hadoop2.7
PYSPARK_DRIVER_PYTHON	Jupyter
PYSPARK_DRIVER_PYTHON_OPTS	notebook

- You also need to set the path where you unpacked your spark to the system's path variable. In the same environment variable settings window, look for the 'Path' variable and click the 'Edit…' button to proceed.

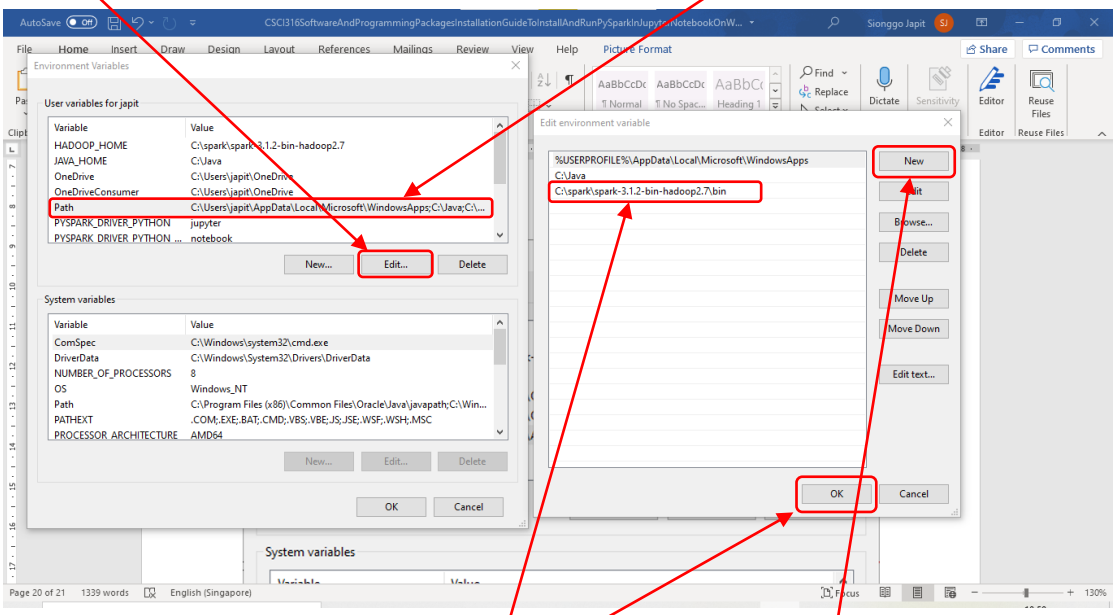

- In the next 'Edit environment variables' pop-up window, click the 'New' button, and enter/add the path where you unpacked your spark packages follows by a '\bin'. For example, I unpacked my spark in C:\spark\spark-3.1.2-bin-hadoop2.7. In this path, there is one directory where the binaries are stored. I will add that directory to the path, hence, my path will be 'C:\spark\spark-3.1.2-bin-hadoop2.7\bin'.

  Click the 'OK' button to save your entry.

- By now, I suppose the installation of Anaconda may have been completed. Check that the Anaconda installation is indeed completed.
- Next, we need install findspark, pyspark, and other required libraries for CSCI316. We will use Anaconda to do so:
  i. Open the Anaconda prompt.

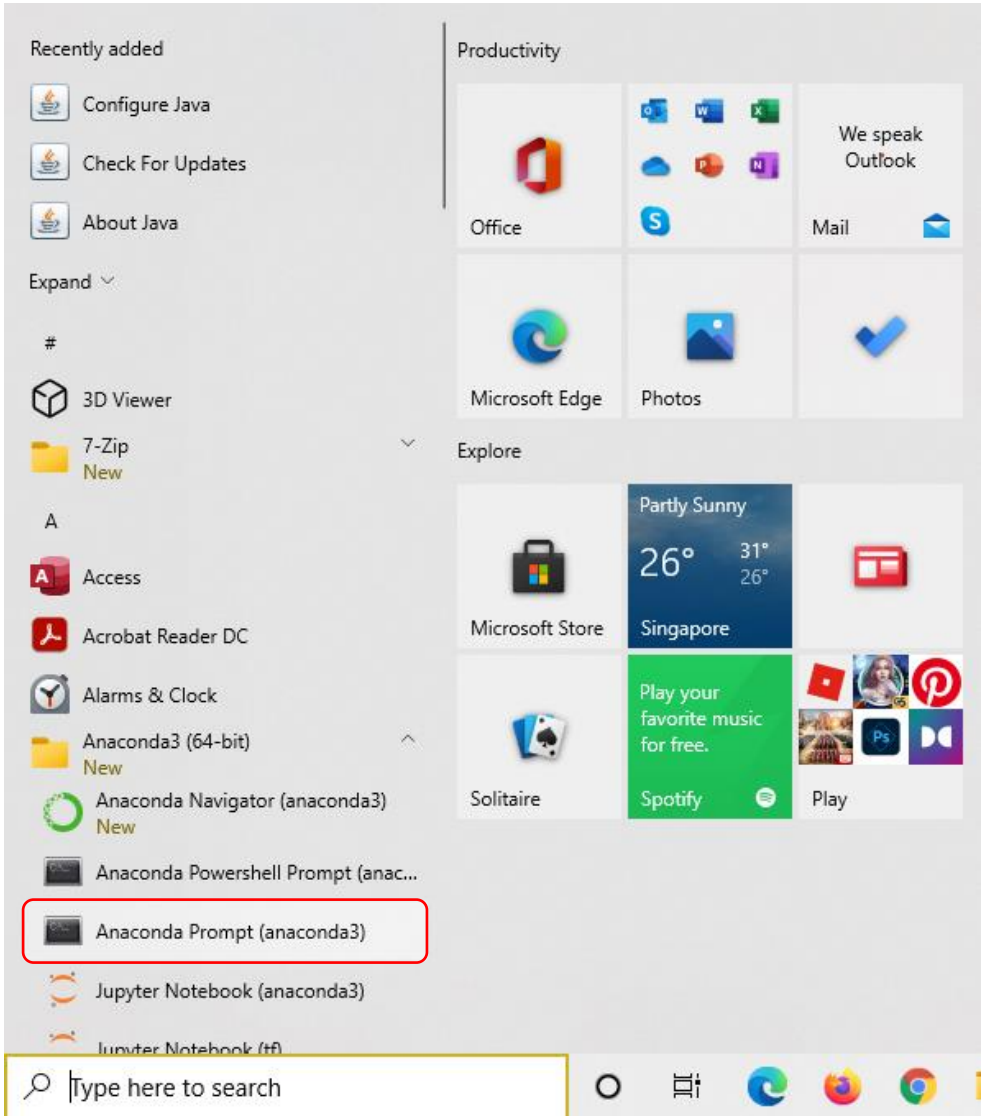

  ii. Installing findspark
    - At the prompt, type **'pip install findspark'** followed with a 'return' key.

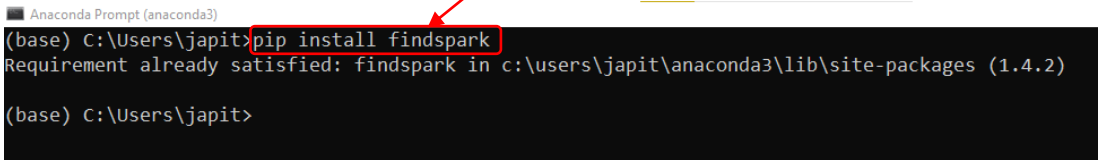

Since I have installed my 'findpark' earlier, the system will indicate that the requirement has already been satisfied, otherwise, the system will install the 'findpark' for you.

iii. Installing pyspark

- At the prompt, type '**pip install pyspark**' followed with a 'return' key.

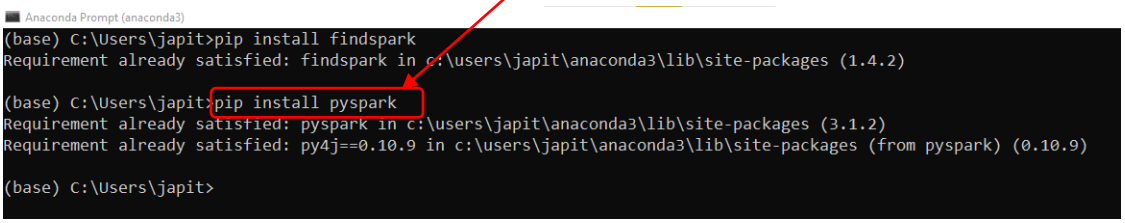

Similarly, since I have installed my 'pyspark' earlier, the system will indicate that the requirement has already been satisfied, otherwise, the system will install the 'pypark' for you.

iv. Installing tensorflow

- At the prompt, type '**pip install –upgrade tensorflow**' followed with a 'return' key.

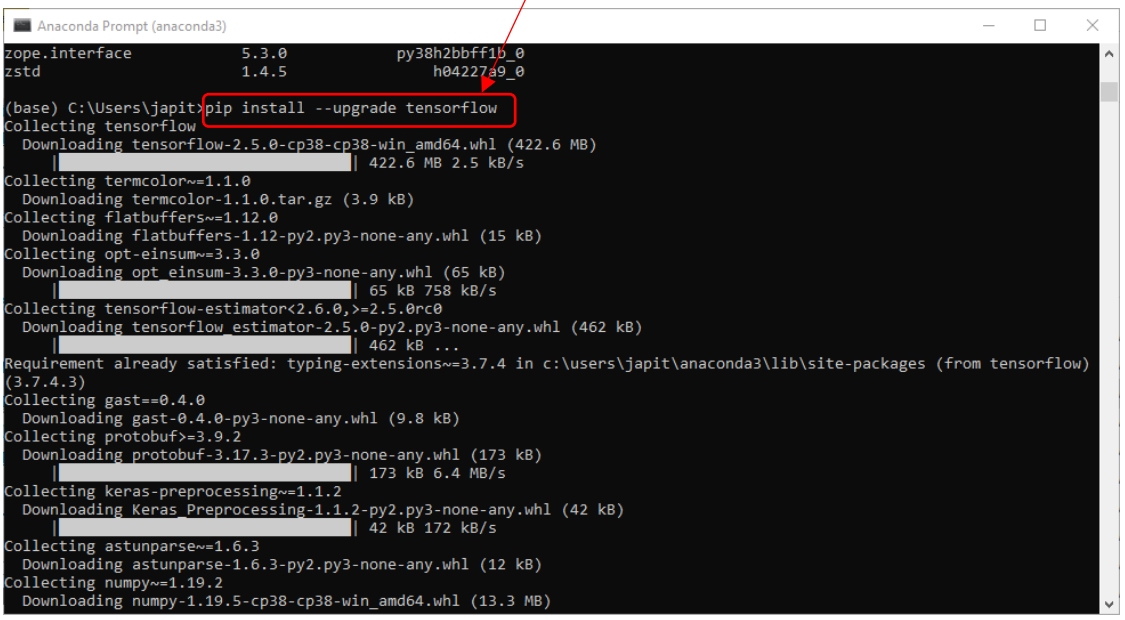

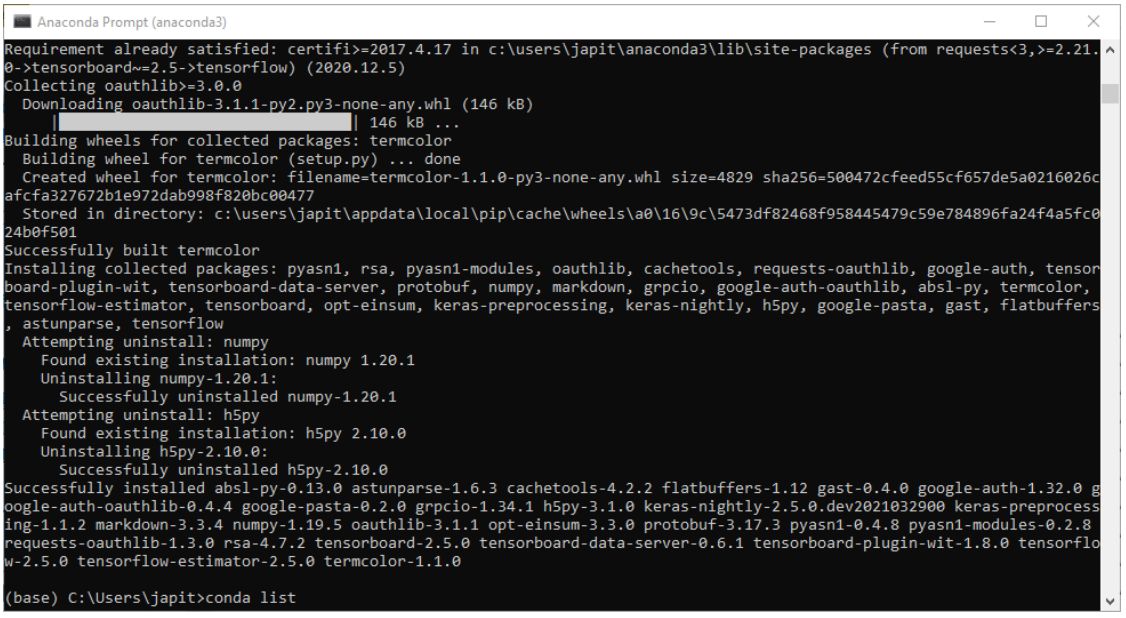

v. Installing pyarrow

- At the prompt, type '**conda install -c conda-forge pyarrow**' followed with a 'return' key.

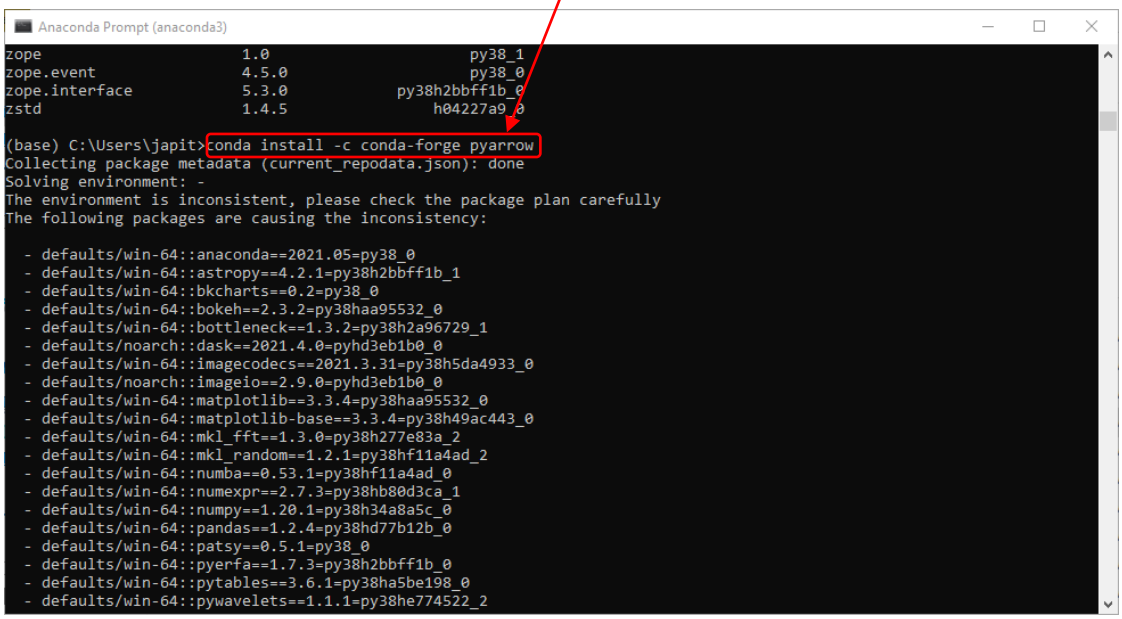

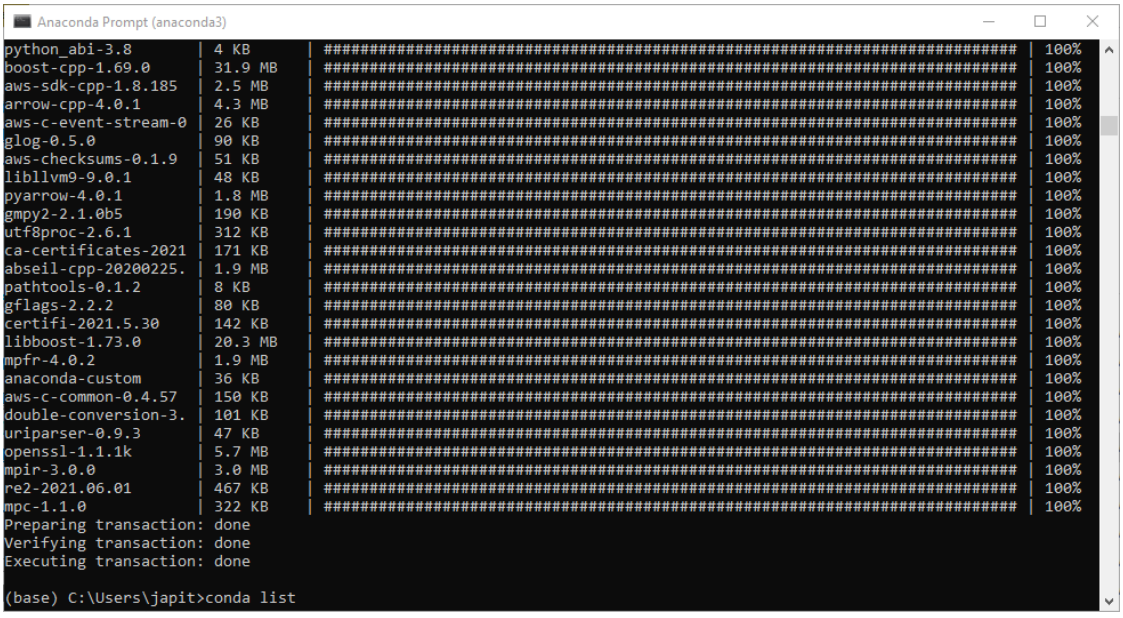

vi. Installing cairocffi

- At the prompt, type '**conda install -c conda-forge/label/cf202003 cairocffi**' followed with a 'return' key.

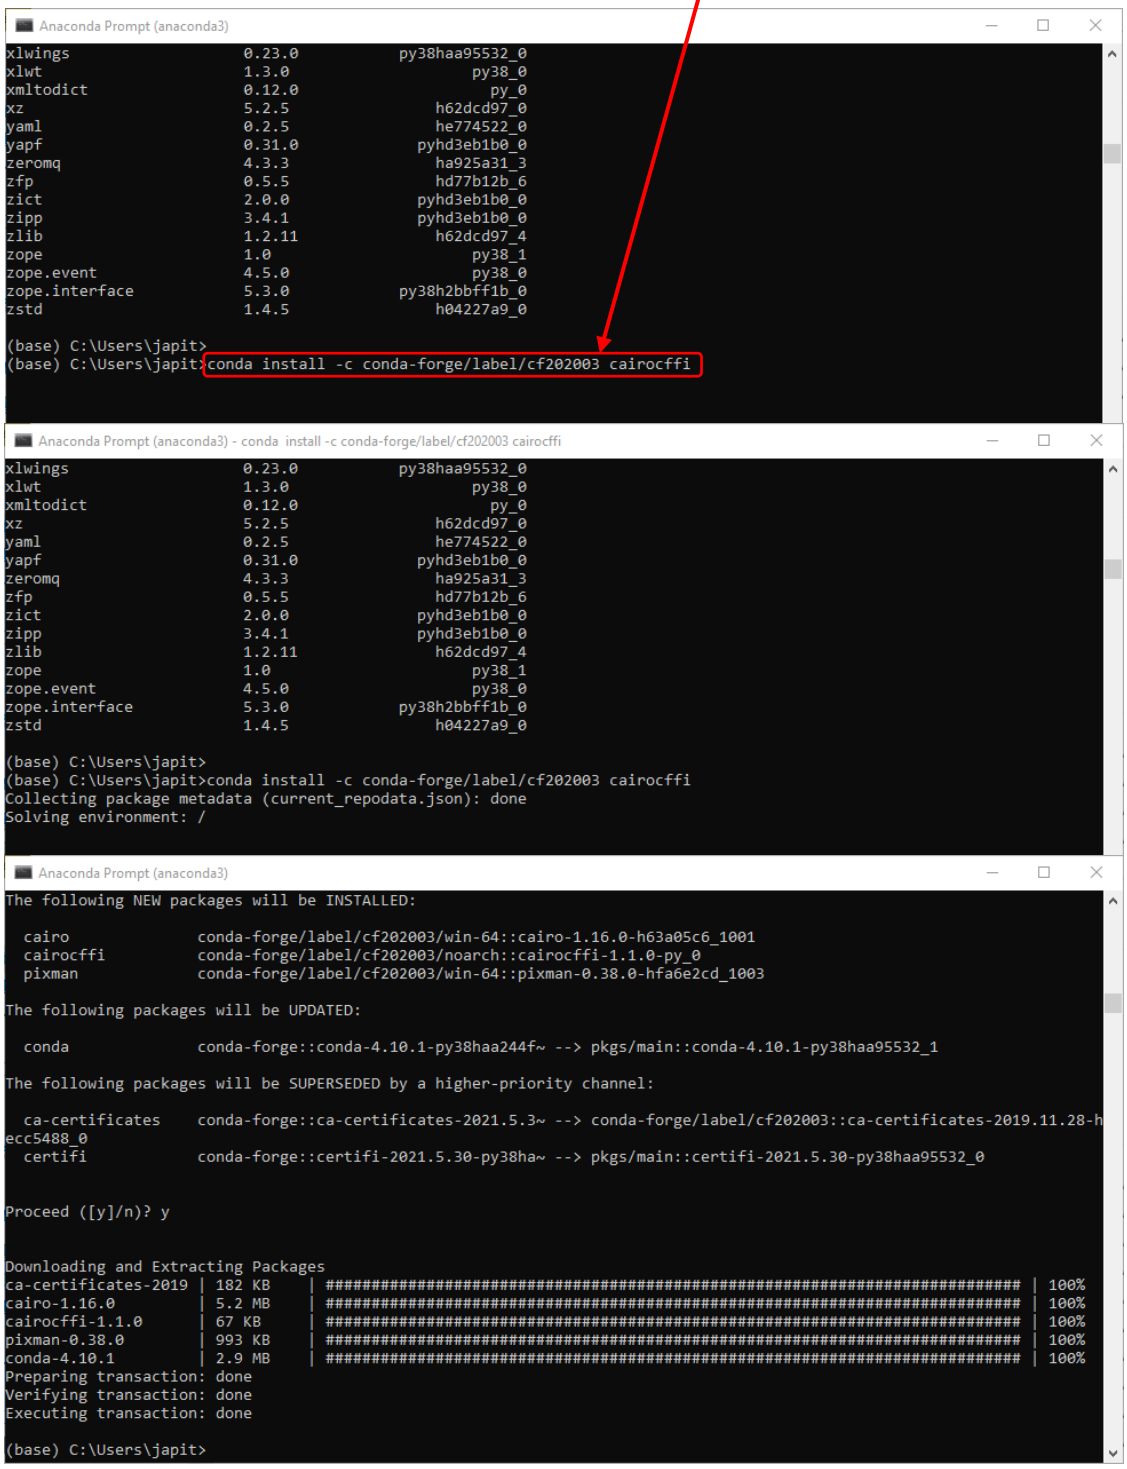

vii. We can now check/verify that all the required software/libraries are installed. To do that at the prompt, type '**conda list**' followed with a 'return' key.

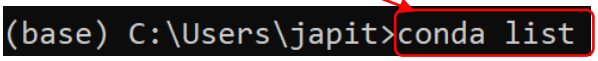

C. Running PySpark in Jupyter Notebook

i. To start Jupyter Notebook, open Jupyter Notebook via the windows '**Start**' icon. Click on the lable '**Jupyter Notebook (anaconda3)**' to start Jupyter Notebook.

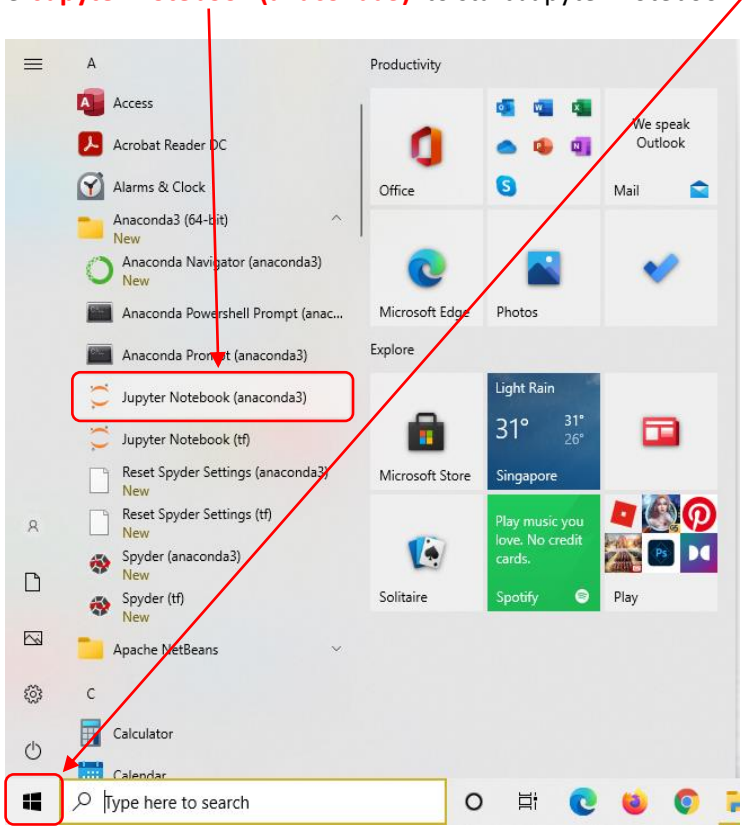

The Jupyter Notebook server is started. Leave this window stays open.

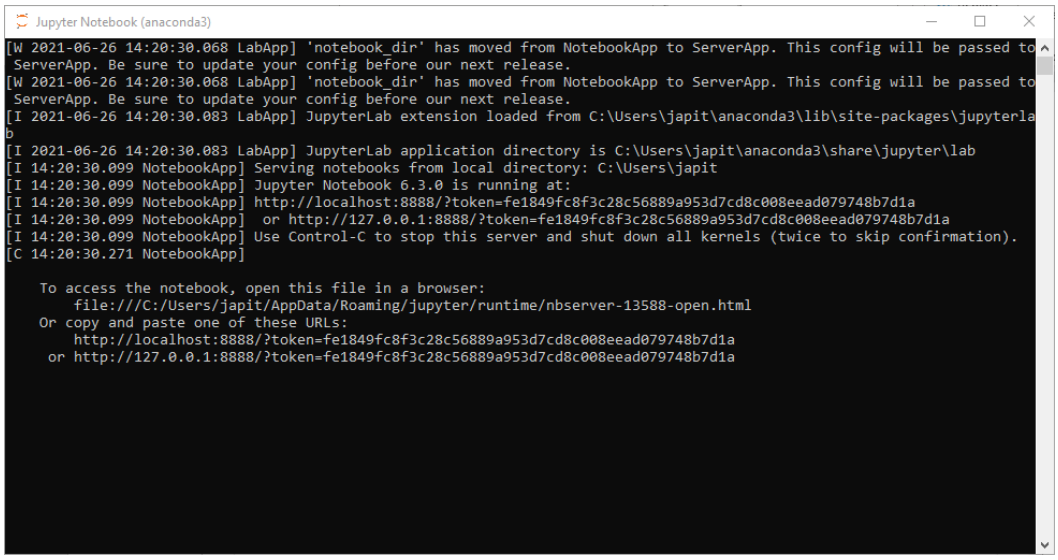

A Jupyter Notebook client is open. Navigate to the desired working directory. In my setup, I have a directory named **'Documents'** created in my user's name in Windows. I will use this directory as my working directory for CSCI316.

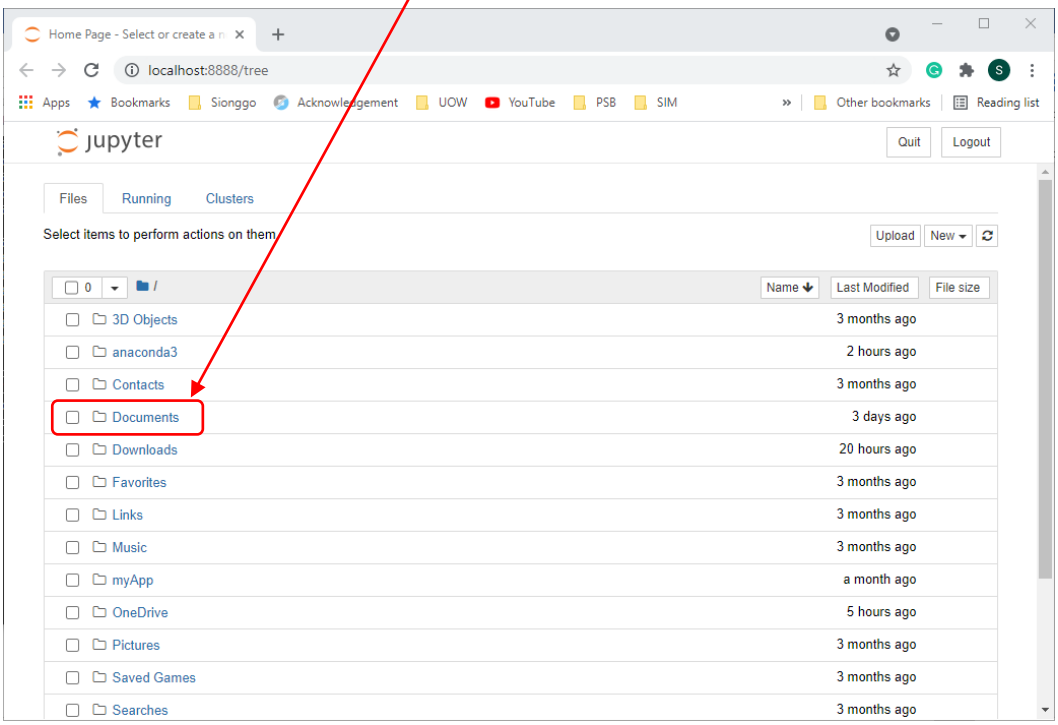

If you do not have a suitable working directory and want to create a new directory, you can click on the New icon and use the option to create a new folder. Once the new folder is created, you can rename it to your choice.

I am now in my working directory. I create a new Python Jupyter Notebook by selecting the '**New**' icon and from the drop-down option, choose '**Python 3**' option.

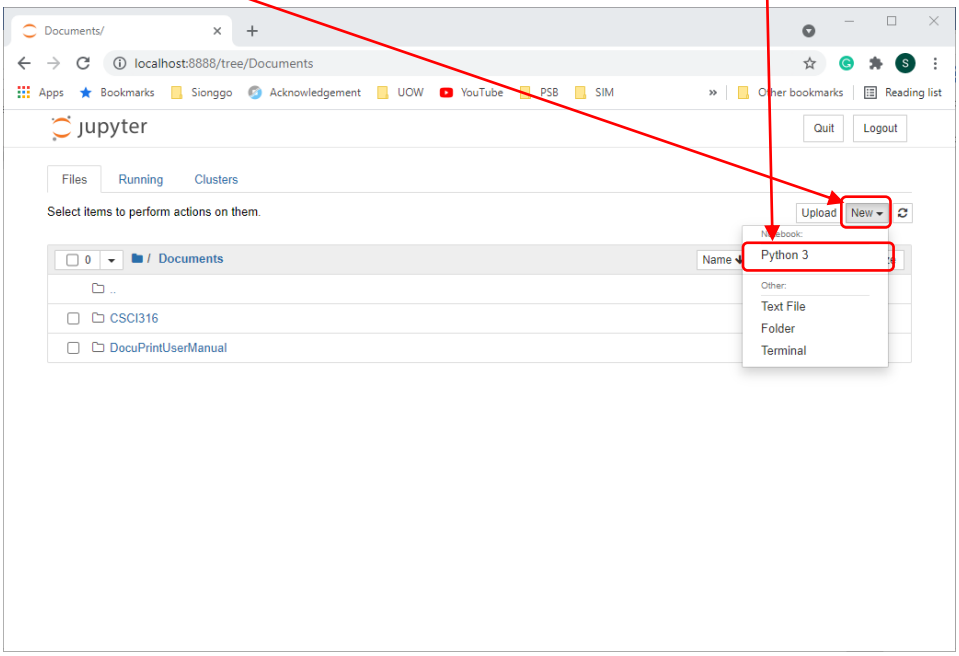

A new Jupyter Notebook node is created.

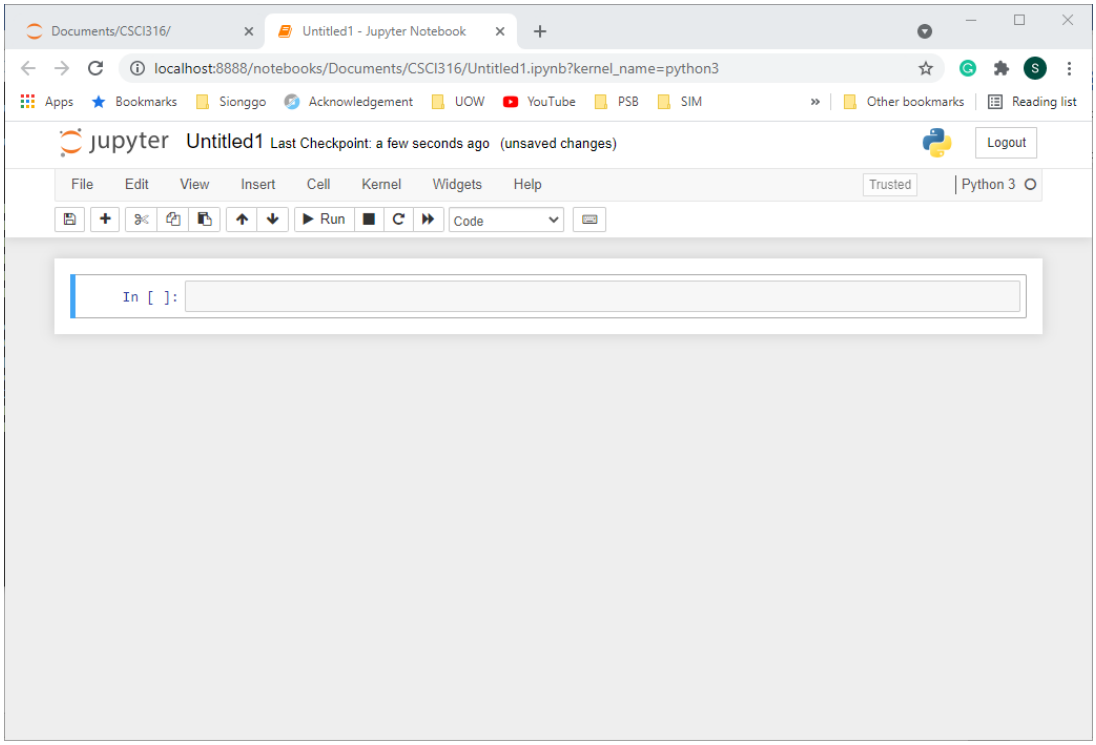

Type the following **codes** in the Jupyter notebook and **run**/execute the code.

```
import findspark
findspark.init()

import pyspark #only run after findspark.init()
from pyspark.sql import SparkSession
spark = SparkSession.builder.getOrCreate()

df = spark.sql("select 'spark' as hello")
df.show()
```

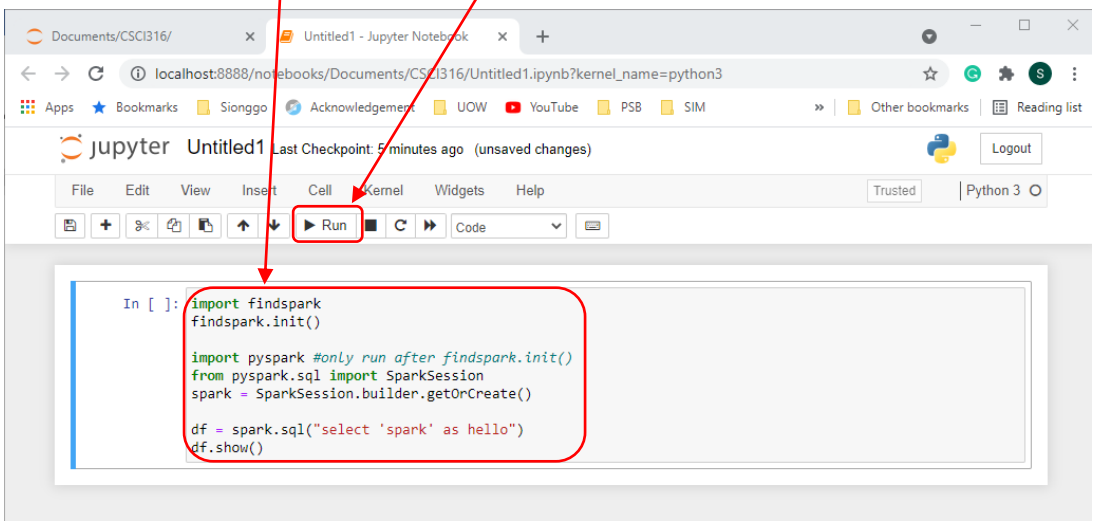

If you see this, ‘Congratulation’ you have successfully install PySpark in your system.