

U

O

W

# A Probabilistic Perspective of Classification

CSCI316 Big Data Mining Techniques and Implementation



UNIVERSITY  
OF WOLLONGONG  
AUSTRALIA

# Contents

Bayes' Theorem

Implementation of simple Naïve Bayes classifier

Key characteristics of NB classifier

# Bayesian Classification

- A probabilistic classifier: performs *probabilistic prediction*, i.e., predicts class membership probabilities
- Foundation: Based on Bayes' Theorem.
- Performance: A simple Bayesian classifier, *naïve Bayesian classifier*, has comparable performance with decision tree and other classifiers
- Incremental: Each training example can incrementally increase/decrease the probability that a hypothesis is correct — prior knowledge can be combined with observed data
- Standard: Even though general Bayesian methods are computationally intractable, simple Bayesian methods can provide a baseline of optimal decision making against which other methods can be measured

# Classification Concepts Recap

- Given a set of records, each of which is described by a sequence of attributes  $X_1, \dots, X_n, Y$ . The last one is an attribute of interest, called a **class**. The rest are called **features**.
- Given a new record where  $Y$  is unknown, the task of classification is to predict which class this record falls into.
- **Probabilistic classifier**: the output of prediction is a class together with a *probabilistic score*
  - to what extent the new record falls into the output class
  - Provides the likelihood instead of a hard decision

# Probability and Uncertainty

- Our main tool is the probability theory, which assigns to a numerical degree of belief between  $0$  and  $1$  to each event.
  - It provides a way of characterizing the uncertainty
- Random variables:
  - Boolean random variables: cavity might be true or false
  - Discrete random variables: weather might be sunny, rainy, cloudy, snow
    - $P(\text{weather} = \text{sunny})$
    - $P(\text{weather} = \text{rainy})$
    - $P(\text{weather} = \text{cloudy})$
    - $P(\text{weather} = \text{snow})$
  - Continuous random variables: the temperature has continuous values
    - Discretization:  $< 10$ ,  $[10, 20]$ ,  $> 20$
    - Probability density function: e.g., Normal distribution.

# Prior and Posterior Probabilities

- Before the evidence is obtained; prior probability
  - $P(a)$  the prior probability that the proposition is true
  - $P(\text{rain}) = 0.1$
- After the evidence is obtained; posterior probability
  - $P(a \mid b)$
  - The probability of  $a$  given that all we know is  $b$  (i.e., **conditional probability**)
  - $P(\text{rain} \mid \text{cloudy}) = 0.8$

# Bayes' Theorem (Simple)

- The conditional probability of event  $C$  occurring, given event  $A$ , is

$$P(C|A) = \frac{P(A \cap C)}{P(A)}$$

- E.g.  $A$  is an attribute and  $C$  is the class.
- Bayes' theorem for two events:
$$P(C|A) = \frac{P(A \cap C)}{P(A)} = \frac{P(C \cap A)}{P(A)} = \frac{P(A|C) \cdot P(C)}{P(A)}$$
  - It links the prior probabilities of two events and their posterior probabilities given each other.

# Example

- Computing the probability that a patient carries a disease based on the result of a lab test.
- The test returns a positive result in 95% of the cases in which the disease is actually present, and it returns a positive result in 6% of the cases in which the disease is not present.
- Furthermore, 1% of the entire population has this disease.
- Let  $C = \{\text{having the disease}\}$  and  $A = \{\text{testing positive}\}$ .
- From the above description,  $P(C) = 0.01$ ,  $P(\neg C) = 0.99$ ,  $P(A|C) = 0.95$  and  $P(A|\neg C) = 0.06$ .



# Reasoning with Bayes' Theorem

$$\begin{aligned}P(A) &= P(A \cap C) + P(A \cap \neg C) \\&= P(C) \cdot P(A|C) + P(\neg C) \cdot P(A|\neg C) \\&= 0.01 \times 0.95 + 0.99 \times 0.06 = 0.0689\end{aligned}$$

$$P(C|A) = \frac{P(A|C)P(C)}{P(A)} = \frac{0.95 \times 0.01}{0.0689} \approx 0.1379$$

Therefore, if some one has a test with positive result, he has 13.79% chance to carry the disease.

# Bayes' Theorem (General)

- In a more general form, Bayes' theorem says that

$$P(Y|X_1, \dots, X_m) = \frac{P(X_1, \dots, X_m|Y) \cdot P(Y)}{P(X_1, \dots, X_m)}$$

- Linking it to classification:  $Y$  is the class and  $X_1, \dots, X_m$  are attributes.
- E.g., attributes: *age, income, student\_status, credit\_rating*; class: *buys\_computer*

<i>RID</i>	<i>age</i>	<i>income</i>	<i>student</i>	<i>credit_rating</i>	<i>Class: buys_computer</i>
1	youth	high	no	fair	no
2	youth	high	no	excellent	no
3	middle_aged	high	no	fair	yes
4	senior	medium	no	fair	yes
5	senior	low	yes	fair	yes

# Naïve Bayes Classifiers

- To apply Bayes theorem to classification, one main problem is *the number of combinations of attribute values*
  - If there are  $m$  attributes and each attribute has  $k$  values, there are  $m^k$  combinations! Impractical to keep track of their joint probabilities.
- Recall  $P(Y|X_1, \dots, X_m) = \frac{P(X_1, \dots, X_m|Y) \cdot P(Y)}{P(X_1, \dots, X_m)}$
- We don't need to compute  $P(X_1, \dots, X_m)$  since we just want to find out *which class (value of  $Y$ ) has the highest score by comparison*.
  - E.g., given *age=youth, income=high, student=no*, and *credit\_rating=fair*, is *buys\_computer=yes* more likely?
    - In this case, we don't need to know the joint probability of *age=youth, income=high, student=no*, and *credit\_rating=fair*
  - In other words, we just reply on
$$P(Y|X_1, \dots, X_m) \propto P(X_1, \dots, X_m|Y) \cdot P(Y)$$
where  $\propto$  indicates “being propositional to”.

# Naïve Bayes Classifiers

- Still,  $P(Y|X_1, \dots, X_m) = \frac{P(X_1, \dots, X_m|Y) \cdot P(Y)}{P(X_1, \dots, X_m)}$
- We use the **conditional independence** assumption.
  - Each attribute is conditionally independent of every other attribute given a class label
  - Namely,  $P(X_1, \dots, X_m|Y) = P(X_1|Y) \cdots P(X_m|Y)$  which dramatically simplifies the computation of  $P(X_1, \dots, X_m|Y)$
- Therefore, we are concerned with

$$P(Y|X_1, \dots, X_m) \propto P(X_1|Y) \cdots P(X_m|Y) \cdot P(Y)$$

# Dataset Example

- Training tuples:

<i>RID</i>	<i>age</i>	<i>income</i>	<i>student</i>	<i>credit_rating</i>	<i>Class: buys_computer</i>
1	youth	high	no	fair	no
2	youth	high	no	excellent	no
3	middle_aged	high	no	fair	yes
4	senior	medium	no	fair	yes
5	senior	low	yes	fair	yes
6	senior	low	yes	excellent	no
7	middle_aged	low	yes	excellent	yes
8	youth	medium	no	fair	no
9	youth	low	yes	fair	yes
10	senior	medium	yes	fair	yes
11	youth	medium	yes	excellent	yes
12	middle_aged	medium	no	excellent	yes
13	middle_aged	high	yes	fair	yes
14	senior	medium	no	excellent	no

# Illustration of Naïve Bayes Classifiers

- Let  $C_1$  correspond to the class *buys\_computer = yes* and  $C_2$  correspond to *buys\_computer = no*.
- Let  $X$  denote  
(*age = youth, income = medium, student = yes, credit rating = fair*)
- The objective is to maximize  $P(X|C_i)P(C_i)$  for  $i = 1, 2$
- First, the prior probability of each class can be computed based on the training tuples:

$$P(\text{buys\_computer} = \text{yes}) = 9/14 = 0.643$$

$$P(\text{buys\_computer} = \text{no}) = 5/14 = 0.357$$

# Illustration of Naïve Bayes Classifiers

- Next, compute the conditional probabilities of attributes on the class labels:

$$P(\text{age} = \text{youth} \mid \text{buys\_computer} = \text{yes}) = 2/9 = 0.222$$

$$P(\text{age} = \text{youth} \mid \text{buys\_computer} = \text{no}) = 3/5 = 0.600$$

$$P(\text{income} = \text{medium} \mid \text{buys\_computer} = \text{yes}) = 4/9 = 0.444$$

$$P(\text{income} = \text{medium} \mid \text{buys\_computer} = \text{no}) = 2/5 = 0.400$$

$$P(\text{student} = \text{yes} \mid \text{buys\_computer} = \text{yes}) = 6/9 = 0.667$$

$$P(\text{student} = \text{yes} \mid \text{buys\_computer} = \text{no}) = 1/5 = 0.200$$

$$P(\text{credit\_rating} = \text{fair} \mid \text{buys\_computer} = \text{yes}) = 6/9 = 0.667$$

$$P(\text{credit\_rating} = \text{fair} \mid \text{buys\_computer} = \text{no}) = 2/5 = 0.400$$

# Naïve Bayes Reasoning

- Using those probabilities, obtain:

$$\begin{aligned}P(X|buys\_computer = yes) &= P(age = youth | buys\_computer = yes) \\&\quad \times P(income = medium | buys\_computer = yes) \\&\quad \times P(student = yes | buys\_computer = yes) \\&\quad \times P(credit\_rating = fair | buys\_computer = yes) \\&= 0.222 \times 0.444 \times 0.667 \times 0.667 = 0.044.\end{aligned}$$

$$P(X|buys\_computer = no) = 0.600 \times 0.400 \times 0.200 \times 0.400 = 0.019.$$

- Finally

$$P(X|buys\_computer = yes)P(buys\_computer = yes) = 0.044 \times 0.643 = 0.028$$

$$P(X|buys\_computer = no)P(buys\_computer = no) = 0.019 \times 0.357 = 0.007$$

- Therefore, the classifier predicts *buys\_computer = yes*



# A Simple NB Implementation

- Assume that our goal is to implement a Naïve Bayes classifier to predict whether a given text contains abusive content or not
  - represented by “1” or “0”, respectively
- In this example, we generate the training dataset by ourselves:

*# the training dataset*

**def** loadDataSet():

```
postingList = [['my', 'dog', 'has', 'flea',  
                'problems', 'help', 'please'],  
               ['maybe', 'not', 'take', 'him',  
                'to', 'dog', 'park', 'stupid'],  
               ['my', 'dalmation', 'is', 'so',  
                'cute', 'I', 'love', 'him'],  
               ['stop', 'posting', 'stupid',  
                'worthless', 'garbage'],  
               ['mr', 'licks', 'ate', 'my', 'steak',  
                'how', 'to', 'stop', 'him'],  
               ['quit', 'buying', 'worthless',  
                'dog', 'food', 'stupid']]
```

```
classVec = [0, 1, 0, 1, 0, 1] # 1 is abusive, 0 not
```

```
return postingList, classVec
```

# A Simple NB Implementation

- We generate a vocabulary of words and a feature matrix where each word is an attribute (i.e., column name)
  - Each row is a vector for a text; if the text contains some word, it has value 1 in the corresponding column.

```
from numpy import *
def createVocabList(dataSet):
    vocabSet = set([]) # create empty set
    for document in dataSet:
        vocabSet = vocabSet | set(document)
        # union of the two sets
    return list(vocabSet)

def setOfWords2Vec(vocabList, inputSet):
    returnVec = [0] * len(vocabList)
    for word in inputSet:
        if word in vocabList:
            returnVec[vocabList.index(word)] = 1
        else:
            print("word: %s is not in my Vocabulary!" % word)
    return returnVec
```

# A Simple NB Implementation

- Check the two functions that we have just defined:

```
# call the two defined functions for illustration:
list0Posts, listClasses = loadDataSet()
myVocabList = createVocabList(list0Posts)
print(myVocabList)
# ['cute', 'love', 'help', 'garbage', 'quit', 'I',
# 'problems', 'is', 'park', 'stop', 'flea',
# 'dalmation', 'licks', 'food', 'not', 'him', 'buying',
# 'posting', 'has', 'worthless', 'ate', 'to', 'maybe',
# 'please', 'dog', 'how', 'stupid', 'so', 'take',
# 'mr', 'steak', 'my']
setOfWords2Vec(myVocabList, list0Posts[0])
# [0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0,
# 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1]
```

- You can generate a training matrix by calling `setOfWords2Vec()` in a loop (for each element in `list0Posts`).

# Pseudo-code of NB Training Function

- The pseudo-code of training function of a Naïve Bayes classifier is as follows:

```
calculate the number (or proportion) of documents in each class;  
for every training document:  
    for each class:  
        if a token appears in the document then  
            increase the count for that token;  
        increase the total count for tokens;  
    for each class:  
        for each token:  
            divide the token count by the total token count  
            to get conditional probabilities;  
return conditional probabilities for each class;
```

- Note that only discrete random variables (e.g., categorical attributes) are considered in this example.
  - Continuous random variables are considered later.

# A Simple NB Implementation

*# build a naive bayes classifier: step 1*

```
def trainNB0(trainMatrix, trainCategory):
    numTrainDocs = len(trainMatrix)
    numWords = len(trainMatrix[0])
    pAbusive = sum(trainCategory) / float(numTrainDocs)
    p0Num = zeros(numWords) # as numerator
    p1Num = zeros(numWords) # as numerator
    p0Denom = 0 # as denominator
    p1Denom = 0 # as denominator
    for i in range(numTrainDocs):
        if trainCategory[i] == 1:
            p1Num += trainMatrix[i]
            p1Denom += sum(trainMatrix[i])
        else:
            p0Num += trainMatrix[i]
            p0Denom += sum(trainMatrix[i])
    p1Vect = p1Num / p1Denom
    p0Vect = p0Num / p0Denom
    return p0Vect, p1Vect, pAbusive
```

# A Simple NB Implementation

```
# build a naive bayes classifier: step 2
# vec2Classify the output vector of setOfWords2Vec, say,
[0, 1, 0,...]
# p0Vec, p1Vec, pAbusive are the output of trainBN0()
def classifyNB0(vec2Classify, p0Vec, p1Vec, pAbusive):
    # element-wise power computation
    p1 = prod(power(p1Vec, vec2Classify)) * pAbusive
    p0 = prod(power(p0Vec, vec2Classify)) * (1.0 -
        pAbusive)
    if p1 > p0:
        return 1
    else:
        return 0
```

Can you see this part is just the Bayesian classifier  
 $P(Y|X_1, \dots, X_m) \propto P(X_1|Y) \cdots P(X_m|Y) \cdot P(Y)$  ?

# A Simple NB Implementation

- Build an NB classifier and test it:

```
def testingNB0():  
    # training  
    list0Posts, listClasses = loadDataSet()  
    myVocabList = createVocabList(list0Posts)  
    trainMat = []  
    for postinDoc in list0Posts:  
        trainMat.append(setOfWords2Vec(myVocabList, postinDoc))  
    p0V, p1V, pAb = trainNB0(array(trainMat),  
                              array(listClasses))  
  
    # classifying: case 1  
    testEntry = ['love', 'my', 'dalmation']  
    thisDoc = array(setOfWords2Vec(myVocabList, testEntry))  
    print(testEntry, 'classified as: ',  
          classifyNB0(thisDoc, p0V, p1V, pAb)) # out: 0  
  
    # classifying: case 2  
    testEntry = ['stupid', 'garbage']  
    thisDoc = array(setOfWords2Vec(myVocabList, testEntry))  
    print(testEntry, 'classified as: ',  
          classifyNB0(thisDoc, p0V, p1V, pAb)) # out: 1
```

# Multiple Occurrences

- In our simple NB implementation, we've treated the presence or absence of a word as a feature.
- But if a word appears more than once in a document, this information is not accounted for.
- Bag-of-words model: a *bag* of words can have multiple occurrences of each word, whereas a *set* of words can have only one occurrence of each word.

```
def bagOfWords2Vec(vocabList, inputSet):  
    returnVec = [0] * len(vocabList)  
    for word in inputSet:  
        if word in vocabList:  
            returnVec[vocabList.index(word)] += 1  
            # was "=" before  
    return returnVec
```



# Numerical Underflow

- If the number of attributes is large, the outputs of a Naïve Bayesian classifier are usually very small.
- In theory this is not a problem, because only the ratio between the outputs matters; however, in practical, the difference may be close or rounded off to 0 (this is unknown as the *underflow* problem).
- To avoid this, one widely used treatment is to manipulate a logarithm of a number rather than the number itself. Therefore,
- Thus  $p_* = p_1 \cdots p_m$  becomes  $\log(p_*) = \log(p_1) + \cdots + \log(p_m)$ 
  - The ratio between the output values of the classifier is not distorted!
  - As multiplication becomes +, the underflow is avoided.

# Smoothing Zero Count

- Another problem is the *zero count*: the count of records with a value of an attribute is zero when some class label is given
- If the zero count occurs, then one of  $P(X_1|Y), \dots, P(X_m|Y)$  is zero, and their multiplication is zero (no matter how large the rest are)
  - This is certainly counter-intuitive
  - Also, applying the log function to a zero probability,  $\log(0)$  is negative infinite
- One common technique to overcome this is the *Laplace smoothing* (or add-one) technique: it adds 1 to all counts.
  - Because usually the training dataset is large (i.e., the total count is large), adding 1 to each count causes minimum effect
  - But if it would cause effect, add a very small number  $\varepsilon > 0$  instead of 1.

# Smoothing Zero Count

- Suppose that for the class *buys computer = yes* in some training database,  $D$ , containing 1000 tuples. We have 0 tuple with *income = low*, 990 tuples with *income = medium*, and 10 tuples with *income = high*.
- Without the Laplacian smoothing, the probabilities of those events are 0, 0.990 (from 990/1000) and 0.010 (from 10/1000), respectively.
- If a tuple has *income = low*, the probability of falling into the class *buys computer = yes* is 0, no matter what values for other attributes!
- With the Laplacian smoothing for the three quantities, adding 1 more tuple for each income value: the probabilities become 0.001 (from 1/1003), 0.988 (from 991/1003) and 0.011 (from 11/1003).
- The above phenomenon won't happen.

# Continuous-Value Features

- We now consider an extension to Naïve Bayesian classifiers which are able to handle continuous-value features.
- If  $X$  is continuous, there are two common approaches to compute  $P(X = a | Y = c)$ :
  - **Discretization/bucketing/binning**: The range of  $X$  is  $(-\infty, a_1], [a_2, b_1], \dots, [a_k, b_{k-1}], [b_k, +\infty)$  for some  $k$ .
  - Assume that  $X$  has a **Gaussian distribution** (a.k.a. normal distribution).
- The following is the *probability density function* (PDF) of a Gaussian distribution with mean  $\mu$  and variance  $\sigma^2$ :

$$f(x, \mu, \sigma) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}},$$

- If we compute the mean value  $\mu_0$  and standard deviation  $\sigma_0$  based on the training data for  $X$  when  $Y = c$ , then  $P(X = x | Y = c)$  *is*  $f(x, \mu_0, \sigma_0)$

# Continuous-Value Features

- Estimation of mean and variance: Given observations  $[x_1, \dots, x_N]$

- mean  $\mu = \frac{1}{N} \sum_{i=1}^N x_i$

- Variance  $\sigma^2 = \frac{1}{N} \sum_{i=1}^N (x_i - \mu)^2 = \left( \frac{1}{N} \sum_{i=1}^N x_i^2 \right) - \mu^2$


- For example, if the incomes are not discretized in the costumer data and are 30, 36, 47, 50, 56, 60, 63, 70, 110 (K dollars) when *buys\_computers* = *yes*, then

- the mean is 58K and
  - the variance is 481.56

- Then

$$P(\text{income} = 47 | \text{buys\_computers} = \text{yes})$$

$$\text{is } \frac{1}{\sqrt{2\pi} \cdot 21.94} e^{-\frac{(47-58)^2}{2 \cdot 481.56}} = 0.016$$




# Continuous-Value Features

- To reason about PDF of the Gaussian distribution, we can use the norm package of the scipy.stats library :

<https://docs.scipy.org/doc/scipy/reference/generated/scipy.stats.norm.html#scipy.stats.norm>

```
import numpy as np
from scipy.stats import norm
import matplotlib.pyplot as plt
a = range(9)
mu = np.mean(a) # mean
sigma = (np.var(a))**0.5
# standard deviation
x = np.arange(-6, 14, 0.1)
y = norm.pdf(x, mu, sigma)
plt.plot(x,y)
plt.show()
```



# \*Continuous-Value Features

- The previous example provides an interpretation is somehow “over simplistic”, since the probability that a continuous random variable takes a particular value is zero.
- Instead, we should compute the conditional probability that  $X$  lies within some interval, say,  $[r, r + \epsilon]$ , where  $\epsilon$  is a small constant:

$$P(r \leq X \leq r + \epsilon) = \int_r^{r+\epsilon} f(X, \mu, \sigma) dX \approx f(X, \mu, \sigma) \cdot \epsilon$$

- Since  $\epsilon$  appears as a constant multiplicative factor for each class, it *cancels out* when normalizing the target probability, leaving just the  $f(X, \mu, \sigma)$  part.

# Comments and Summary

- Advantages
  - Easy to implement
  - Good results obtained in most of the cases
- Disadvantages
  - Assumption: class conditional independence, therefore loss of accuracy
  - Practically, dependencies exist among variables
    - E.g., Symptoms: fever, cough etc., Disease: lung cancer, diabetes, etc.
    - Dependencies among these cannot be modeled by Naïve Bayes Classifier
- How to deal with these dependencies?
  - Belief Bayesian Network
- From correlation to causality?